



## Lab 9      Sorting

### Goal

In this lab the performance of basic sorting algorithms will be explored.

### Resources

- Chapter 8: An Introduction to Sorting

### Java Files

- `SortArray.java`
- `SortDriver.java`

The basic sorts have already been implemented in the `SortArray` class. You will make a new class `SortArrayInstrumented` that will be based on that class. It will allow you to gather statistics about the sorts. The `SortDriver` class will generate the arrays, call the sorts, and then display the statistical results.

### Adding Statistics to Selection Sort

**Step 1.** If you have not done so, look at the implementation of the sorts in `SortArray.java`. Look at the skeleton in `SortDriver.java`. Compile the classes `SortArray`, and `SortDriver`. Run the main method in `SortDriver`.

*Checkpoint: The program will ask you for an array size. Enter 20. An array of 20 random values between 0 and 20 should be generated and displayed. Selection sort will be applied to array and the sorted array will be displayed. Verify that this works correctly.*

*The first goal is to create a new class `SortArrayInstrumented` that will be used to collect statistics about the performance of the sorts. Private variables of the class will be used to record the number of comparisons made.*

**Step 2.** Create a new class name `SortArrayInstrumented`.

**Step 3.** Copy the contents of `SortArray` into `SortArrayInstrumented`. Change the name in the class declaration from `SortArray` to `SortArrayInstrumented`.

**Step 4.** Create a default constructor that does nothing. (It will have work to do later.)

**Step 5.** Remove `static` from all the methods in the `SortArrayInstrumented` class.

*Checkpoint: You should be able to compile `SortArrayInstrumented` without errors.*

*Since the sort methods are no longer static, `SortDriver` must be changed to create an instance of `SortArrayInstrumented` and then invoke the sort method using the instance.*

**Step 6.** In main of `sortDriver` declare and create a new instance of `SortArrayInstrumented` named `sai`.



**Step 7.** Change `SortArray.selectionSort(data, arraySize)` to `sai.selectionSort(data, arraySize)`.

*Checkpoint: Compile and run the program. Enter 20 for the array size. Verify that this works correctly.*

*The next goal is to add code to the selection sort to count the number of times that a comparison of data values is made. Methods will be added to the `SortArrayInstrumented` class to allow the number of comparisons to be recovered.*

**Step 8.** Add a private variable `comparisons` of type `long` to the `SortArrayInstrumented` class. Initialize it to zero in the constructor.

**Step 9.** Add a public accessor method `getComparisons` to the `SortArrayInstrumented` class.

**Step 10.** In order to count the number of times that `compareTo()` is called by selection sort, put the line `comparisons++;` just before the `if` statement in `indexOfSmallest()`. If the code is inserted inside the `then` clause, only the comparisons that result in `true` will be counted.

**Step 11.** In `SortDriver`, add the line `System.out.println(" comparison made: "+sai.getComparisons());` after the call to `selection sort`.

*Checkpoint: Compile and run the program. Enter 20 for the array size. Verify that the sort still works correctly. The number of comparisons should be 190.*

*The next goal is to compute the average number of comparisons made by the sort with many different arrays (all of the same size). Only `SortDriver` will be changed.*

**Step 12.** In `SortDriver`, use the method `getInt()` to set the variable `trials`.

**Step 13.** Starting with the call to `generateRandomArray`, wrap the remainder of the code in `main` in `SortDriver` with a `for` loop that runs the given number of trials.

*Checkpoint: Compile and run the program. Enter 20 for the array size. Enter 3 for the number of trials. Verify that each of the three trials sorted the values correctly and is for a different array of 20 values. The number of comparisons should be 190, 380, and 570.*

*Notice that the number of comparisons gives a running total for all calls. The next goal is to compute and report the minimum and maximum number of comparisons made over all the calls to the sort. To do this, the use of the `comparisons` variable will be changed slightly. It will only be the number of comparisons made by the last call to the sort. The total number of comparisons made by all calls will be held in a new variable. This aids in the computation of the maximum and minimum.*

**Step 14.** Add a private variable `totalComparisons` of type `long` to the `SortArrayInstrumented` class. Initialize it to zero in the constructor.

**Step 15.** Add a private variable `minComparisons` of type `long` to the `SortArrayInstrumented` class. Initialize it to `Long.MAX_VALUE` in the constructor.



**Step 16.** Add a private variable `maxComparisons` of type `long` to the `SortArrayInstrumented` class. Initialize it to zero in the constructor.

**Step 17.** Add three public accessor methods (one for each of the new variables) to the `SortArrayInstrumented` class.

*To compute the minimum and maximum number of comparisons, code needs to be added at the beginning and end of the sort. While the needed code could be added directly to the sorts, it is better to encapsulate it in a couple new methods.*

**Step 18.** Add a private method `startStatistics()` to the `SortArrayInstrumented` class. It should initialize `comparisons` to zero.

**Step 19.** Add a private method `endStatistics()` to the `SortArrayInstrumented` class. It should add `comparisons` to `totalComparisons`. It should compare `comparisons` to `minComparisons` and set `minComparisons` to whichever is smaller. It should also set `maxComparisons` in an analogous fashion.

**Step 20.** Call `startStatistics()` at the beginning of the `selectionSort` method. Call `endStatistics()` at the end of the `selectionSort` method.

**Step 21.** After the `for` loop in `main` of `SortDriver`, add in three statements that print the total, minimum, and maximum number of comparisons.

*Checkpoint: Compile and run the program. Enter 20 for the array size. Enter 3 for the number of trials. Verify that each of the three trials sorted the values correctly and is for a different array of 20 values. The number of comparisons should be 190 for each of the three trials. The total should be 570 and the minimum and maximum should both be 190. Refer to the pre-lab exercises and compare.*

*Enter 10 for the array size. Enter 3 for the number of trials. Verify that each of the three trials sorted the values correctly and is for a different array of 10 values. The number of comparisons should be 45 for each of the three calls. The total should be 135 and the minimum and maximum should both be 45.*

**Step 22.** Compute the average number of comparisons made over the trials and print it. (The average is the total number of comparisons divided by the number of trials.)

**Step 23.** In preparation for filling in the table, comment out the print statements inside the `for` loop in `main`.

*Final checkpoint: Compile and run the program. Enter 20 for the array size. Enter 1000 for the number of trials. The total should be 19000 and the average, minimum, and maximum should all be 190.*

**Step 24.** Fill in this table and record the average in the appropriate column in the table at the end of the directed lab. Use 100 trials.



## Comparisons for Selection Sort

	MINIMUM COMPARISONS	AVERAGE COMPARISONS	MAXIMUM COMPARISONS
Size=10			
Size=50			
Size=100			
Size=200			
Size=300			
Size=400			
Size=500			
Size=750			
Size=1000			

## Adding Statistics to Recursive Insertion Sort

*Most of the work needed has been done before. It is now just a matter of adding the appropriate code to the insertion sort code.*

**Step 25.** Add calls to `startStatistics()` and `endStatistics()` to the public, nonrecursive `insertionSort()` method.

**Step 26.** In the `insertInOrder()` method place code to add one to `comparisons` when `compareTo()` is invoked.

**Step 27.** In main in `SortDriver`, change the call from `selectionSort` to `insertionSort`.

**Step 28.** Uncomment the `print` statements in the `for` loop in main in `SortDriver`.

*Checkpoint: Compile and run the program. Enter 20 for the array size. Enter 3 for the number of trials. Verify that each of the three trials sorted the values correctly and is for a different array of 20 values. The number of comparisons should typically be in the range of 85 to 130 for each of the three with an average of approximately 107. If you get values that are outside this range, retry the test a few times. If you consistently get results outside the range, check the code you added for errors. Verify that the total, minimum, and maximum are correct for the reported number of comparisons.*

*Enter 10 for the array size. Enter 3 for the number of trials. Verify that each of the three trials sorted the values correctly and is for a different array of 10 values. The number of comparisons should be approximately 28 for each of the three trials. Verify that the total, minimum, and maximum are correct for the reported number of comparisons.*

**Step 29.** Uncomment the `print` statements from the previous step.

*Final checkpoint: Compile and run the program. Enter 20 for the array size. Enter 10000 for the number of trials. The average should be approximately 107.*

**Step 30.** Fill in this table and record the average in the appropriate column in the table at the end of the directed lab. Use 100 trials.

*Warning: Depending on the computer you are using, there may be a limit on the number of recursive calls that can be made. If this happens you will get the error `java.lang.StackOverflowError`. While this often indicates that you have entered an infinite recursion, as long as the sort worked for a smaller array, that is not the case here.*

*Adapted from Dr. Hoot's Lab Manual for Data Structures and Abstractions with Java™*



If we examine the pattern of calls, we see that we need to make one recursive call for each entry in the array. As you apply the algorithm to larger and larger arrays, the size of the stack must be larger and we can hit the upper limit. If this happens, find the limit and mark it on the table.

## Comparisons for Insertion Sort

	MINIMUM COMPARISONS	AVERAGE COMPARISONS	MAXIMUM COMPARISONS
Size=10			
Size=50			
Size=100			
Size=200			
Size=300			
Size=400			
Size=500			
Size=750			
Size=1000			

## Average Comparisons for the Two Sorts

	SELECTION SORT	INSERTION SORT
Size=10		
Size=50		
Size=100		
Size=200		
Size=300		
Size=400		
Size=500		
Size=750		
Size=1000		

## Post-Lab Follow-Ups

1. Add a `reset()` method to `SortArrayInstrumented` that will set each of the variables as the constructor does. Modify `SortDriver` to compute the average, minimum, and maximum for each of the two sorts with the input array size and number of trials.
2. Add variables and methods to `SortArrayInstrumented` to compute the total of the squares of the number of comparisons. If the number of comparisons made by three calls were 3, 5, and 2, the sum of the squares would be  $9 + 25 + 4$ . The variance of  $k$  values is the average of the squares of the values minus the square of the average. For the given values, the average is  $10/3$  and the variance is  $38/3 - (10/3)^2$ . The standard deviation is the square root of the variance. Use this to compute and display the standard deviation in `SortDriver`.
3. Another way of measuring the performance of a sort is by the amount of data movement it must do. Anytime an assignment is made using the array, add one to the number of moves. For example, a swap operation would add 3 to the number of moves. Add variables and methods to `SortArrayInstrumented` to compute the total, minimum, and maximum number of moves. Add code to `SortDriver` to display them.



**Note:** This measure is relatively unimportant in Java since the sorts work with arrays of references to objects. Because of this, only references are being moved and not the objects themselves. The time to complete a swap in Java will not depend on the size of object.

- Bubble sort is an older sort whose performance is not competitive with the other basic sorts. Outside of this exercise, you should not use bubble sort. One variant of bubble sort that works on the first  $n$  entries in an array uses the following algorithm.

```
Set first position to 0
Set last position to  $n-2$ 
While the first position is less than the last position
    Set last swap location to first position
    Loop  $i$  from first position to last position
        If the entries at positions  $i$  and  $i+1$  are out of order
            Swap the entries in positions  $i$  and  $i+1$ 
            Set last swap to  $i$ 
Set last position to last swap
```

Implement this version of bubble sort and add statistics to it in `SortArrayInstrumented`. Compute the minimum, maximum, and average number of comparisons done.

- The given Shell sort works with increments of 1, 2, 4, 8, ... (written in reverse order). The performance of Shell's sort can be improved by allowing more varied overlap between the sequences. Implement two new versions of Shell's sort that use the sequence of increments

1, 4, 13, 40, 121      ( $s_i = 3s_{i-1} + 1$ )

and

1, 3, 7, 15, 31      ( $s_i = 2s_{i-1} + 1$ )

Compute the minimum, maximum, and average number of comparisons done by these two versions of Shell's sort and compare with the original.

- Create a new sort that works in two phases. In the first phase, generate  $n \log_2 n$  pairs of random positions in the array that are at least  $(\log_2 n)^2$  apart. If the items in the pair of positions are out of order, switch them. In the second phase, do an insertion sort. Compute the minimum, maximum and average number of comparisons done by this sort and compare with the standard insertion sort. (The idea behind this sort is that insertion sort works much better on arrays that are nearly sorted. In the first phase we do a small amount of work to get the array closer to being in sorted order. Swapping values that are farther apart has a potentially better improvement so we don't want the positions to be too close to each other. Shell sort uses the same idea, but in a more regimented fashion with a decreasing distance between positions.
- By counting the number of inversions in an array, you get a measure of how close the array is to being sorted. Consider every pair of values. (There are  $n(n-1)/2$  pairs.) Each pair that is out of order contributes one to the number of inversions. Implement a method that counts the number of inversions in an array of `Comparable` objects.
- While measuring the performance of a sort against randomly generated arrays is important, in real life data are often not random. Nearly sorted data are frequently encountered. Develop a method that generates random arrays that have at most  $k$  inversions. Use this method to compute the performance of the three sorts for  $k = n/2$ ,  $k = n$ , and  $k = 2n$ .