



**CS/COE 1501**  
**Algorithm Implementation**

# Course Introduction Algorithm Analysis Exhaustive Search

Fall 2018

Sherif Khattab

[ksm73@pitt.edu](mailto:ksm73@pitt.edu)

6307 Sennott Square

(Slides are adapted from Dr. Ramirez's and Dr. Farnan's CS1501 slides.)

# Contact Info

- **Course website:**  
<http://www.cs.pitt.edu/~skhattab/cs1501/>
- **Instructor:** Sherif Khattab [ksm73@pitt.edu](mailto:ksm73@pitt.edu)
  - OH: MW 10-4pm (other times available by appointment)
    - Please book at <http://khattab.youcanbook.me>
  - Office: 6307 Sennott Square
- **Teaching Assistants**
  - Ziyu Zhang for Section 1270 and 1280
  - Hanzhong Zheng and Jie Liu (Grader) for Section 1400 and 1410
  - No recitations this week

# Textbook



- **Algorithms (4<sup>th</sup> Edition)**
- Robert Sedgewick and Kevin Wayne

# Grades

- 40% exams: 26% on higher grade and 14% on lower
- 40% on best four out of five assignments worth 10% each
  - posted and submitted on **Courseweb**
- 10% recitation participation
- 10% mini-quizzes: on Top Hat during each lecture.  
The join code for Top Hat is 064642.

Up until now, your classes have focused on how you *could* solve a problem. Here, we will start to look at how you *should* solve a problem.

# Alright, then, down to business...

First some definitions:

- **Offline problem**
  - We provide the computer with some input and after some time receive some acceptable output
- **Algorithm**
  - A step-by-step procedure for solving a problem or accomplishing some end
- **Program**
  - An algorithm expressed in a language the computer can understand

An algorithm solves a problem if it produces an acceptable output on *every* input

# Goals of the course (1)

- To learn to convert non-trivial *algorithms* into *programs*
  - Many seemingly simple algorithms can become much more complicated as they are converted into programs
  - Algorithms can also be very complex to begin with, and their implementation must be considered carefully
  - Various issues will always pop up during *implementation*
    - Such as?...

# Example

- Pseudocode for dynamic programming algorithm for relational query optimization
- The optimizer portion of the PostgreSQL codebase is over 28,000 lines of code (i.e., not counting blank/comment lines)

```
Input: SPJ query  $q$  on relations  $R_1, \dots, R_n$ 
Output: A query plan for  $q$ 

1: for  $i = 1$  to  $n$  do {
2:     optPlan( $\{R_i\}$ ) = accessPlans( $R_i$ )
3:     prunePlans(optPlan( $\{R_i\}$ ))
4: }
5: for  $i = 2$  to  $n$  do {
6:     for all  $S \subset \{R_1, \dots, R_n\}$  such that  $|S| = i$  do {
7:         optPlan( $S$ ) =  $\emptyset$ 
8:         for all  $O \subset S$  do {
9:             optPlan( $S$ ) = optPlan( $S$ )  $\cup$  joinPlans(optPlan( $O$ ), optPlan( $S \setminus O$ ))
10:            prunePlans(optPlan( $S$ ))
11:        }
12:    }
13: }
14: finalizePlans(optPlan( $\{R_1, \dots, R_n\}$ ))
15: prunePlans(optPlan( $\{R_1, \dots, R_n\}$ ))
16: return optPlan( $\{R_1, \dots, R_n\}$ )
```

Fig. 1. “Classic” dynamic programming algorithm.

# Goals of the course (2)

- To see and understand differences in algorithms and how they affect the run-times of the associated programs
  - Different algorithms can be used to solve the same problem
  - Different solutions can be compared using many metrics
    - Run-time is a big one
      - Better run-times can make an algorithm more desirable
      - Better run-times can sometimes make a problem solution feasible where it was not feasible before
    - There are other metrics, though...

# How to determine an algorithm's performance

- Implement it and measure performance
  - Any problems with this approach?
- Algorithm Analysis
  - Determine *resource usage* as a function of *input size*
  - Measure *asymptotic* performance
    - Performance as input size increases to infinity

# Let's consider ThreeSum example from text

- Problem:
  - Given a set of arbitrary integers (could be negative), find out how many distinct triples sum to exactly zero
  - Simple solution: triple for loops!

```
public static int count(int[] a) {  
    int n = a.length;  
    int cnt = 0;  
    for (int i = 0; i < n; i++) {  
        for (int j = i+1; j < n; j++) {  
            for (int k = j+1; k < n; k++) {  
                if (a[i] + a[j] + a[k] == 0) {  
                    cnt++;  
                }  
            }  
        }  
    }  
    return cnt;  
}
```

# Definition of Big O?

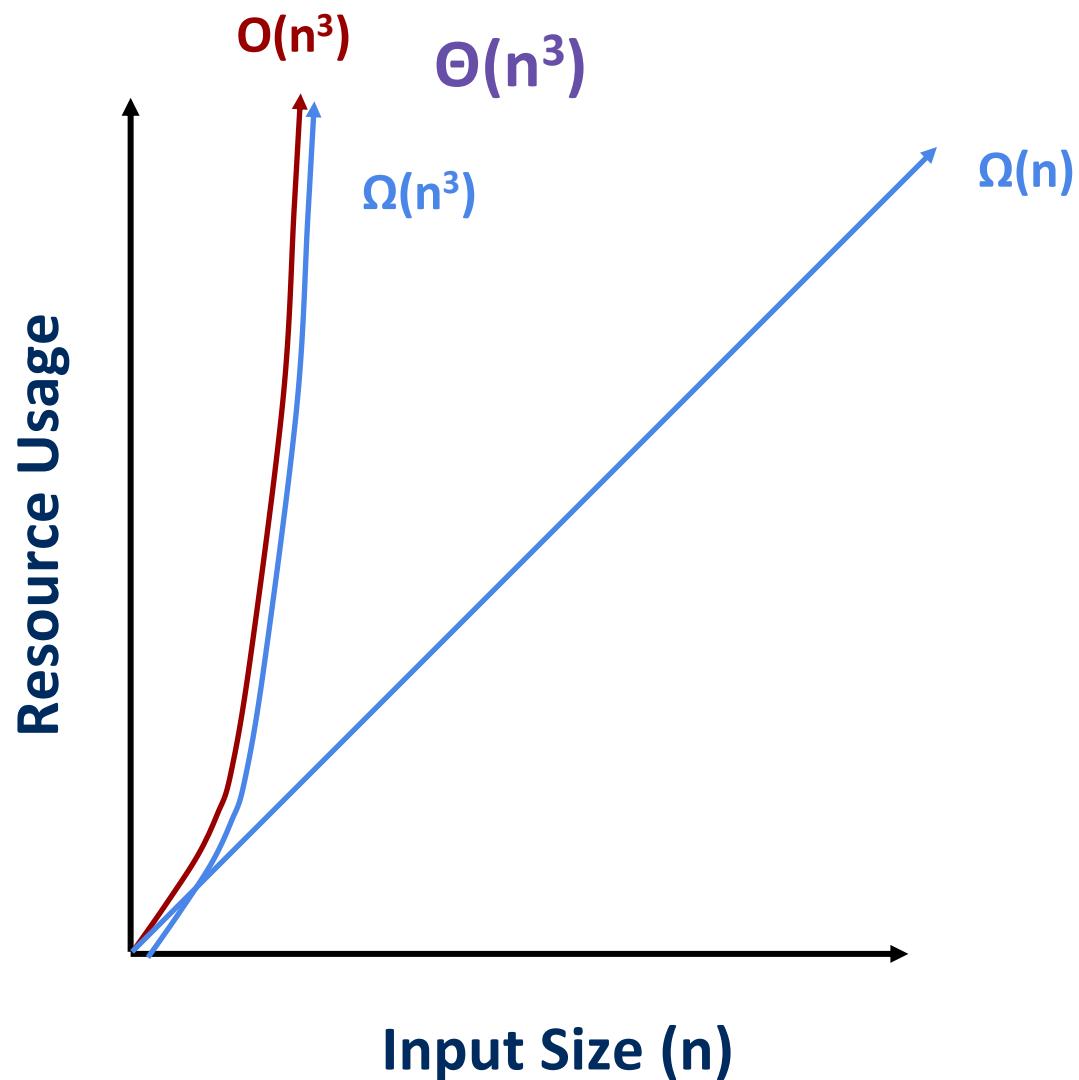
- Big O
  - Upper bound on asymptotic performance
    - As we go to infinity, function representing resource consumption will not exceed specified function
      - E.g., Saying runtime is  $O(n^3)$  means that as input size ( $n$ ) approaches infinity, actual runtime will not exceed  $n^3$

# Wait...

- Assuming that definition...
  - Is ThreeSum  $O(n^4)$ ?
  - What about  $O(n^5)$ ?
  - What about  $O(3^n)??$
- If all of these are true, why was  $O(n^3)$  what we jumped to to start?

# Big O isn't the whole story

- Big Omega
  - Lower bound on asymptotic performance
- Theta
  - Upper and Lower bound on asymptotic performance
  - Exact bound



# Formal definitions

- $f(x)$  is  $O(g(x))$  if constants  $c$  and  $x_0$  exist such that:
  - $|f(x)| \leq c * |g(x)| \forall x > x_0$
- $f(x)$  is  $\Omega(g(x))$  if constants  $c$  and  $x_0$  exist such that:
  - $|f(x)| \geq c * |g(x)| \forall x > x_0$
- if  $f(x)$  is  $O(g(x))$  and  $\Omega(g(x))$ , then  $f(x)$  is  $\Theta(g(x))$ 
  - $c_1, c_2$ , and  $x_0$  exist such that:
    - $c_1 * |g(x)| \leq |f(x)| \leq c_2 * |g(x)| \forall x > x_0$
- May also see  $f(x) \in O(g(x))$  or  $f(x) = O(g(x))$  used to mean that  $f(x)$  is  $O(g(x))$ 
  - Same for  $\Omega$  and  $\Theta$

# Mathematically modelling runtime

- Runtime primarily determined by two factors:
  - Cost of executing each statement
    - Determined by machine used, environment running on the machine
  - Frequency of execution of each statement
    - Determined by program and input

# Let's consider ThreeSum example from text

```
public static int count(int[] a) {  
    int n = a.length;  
    int cnt = 0;  
    for (int i = 0; i < n; i++) {  
        for (int j = i+1; j < n; j++) {  
            for (int k = j+1; k < n; k++) {  
                if (a[i] + a[j] + a[k] == 0) {  
                    cnt++;  
                }  
            }  
        }  
    }  
    return cnt;  
}
```

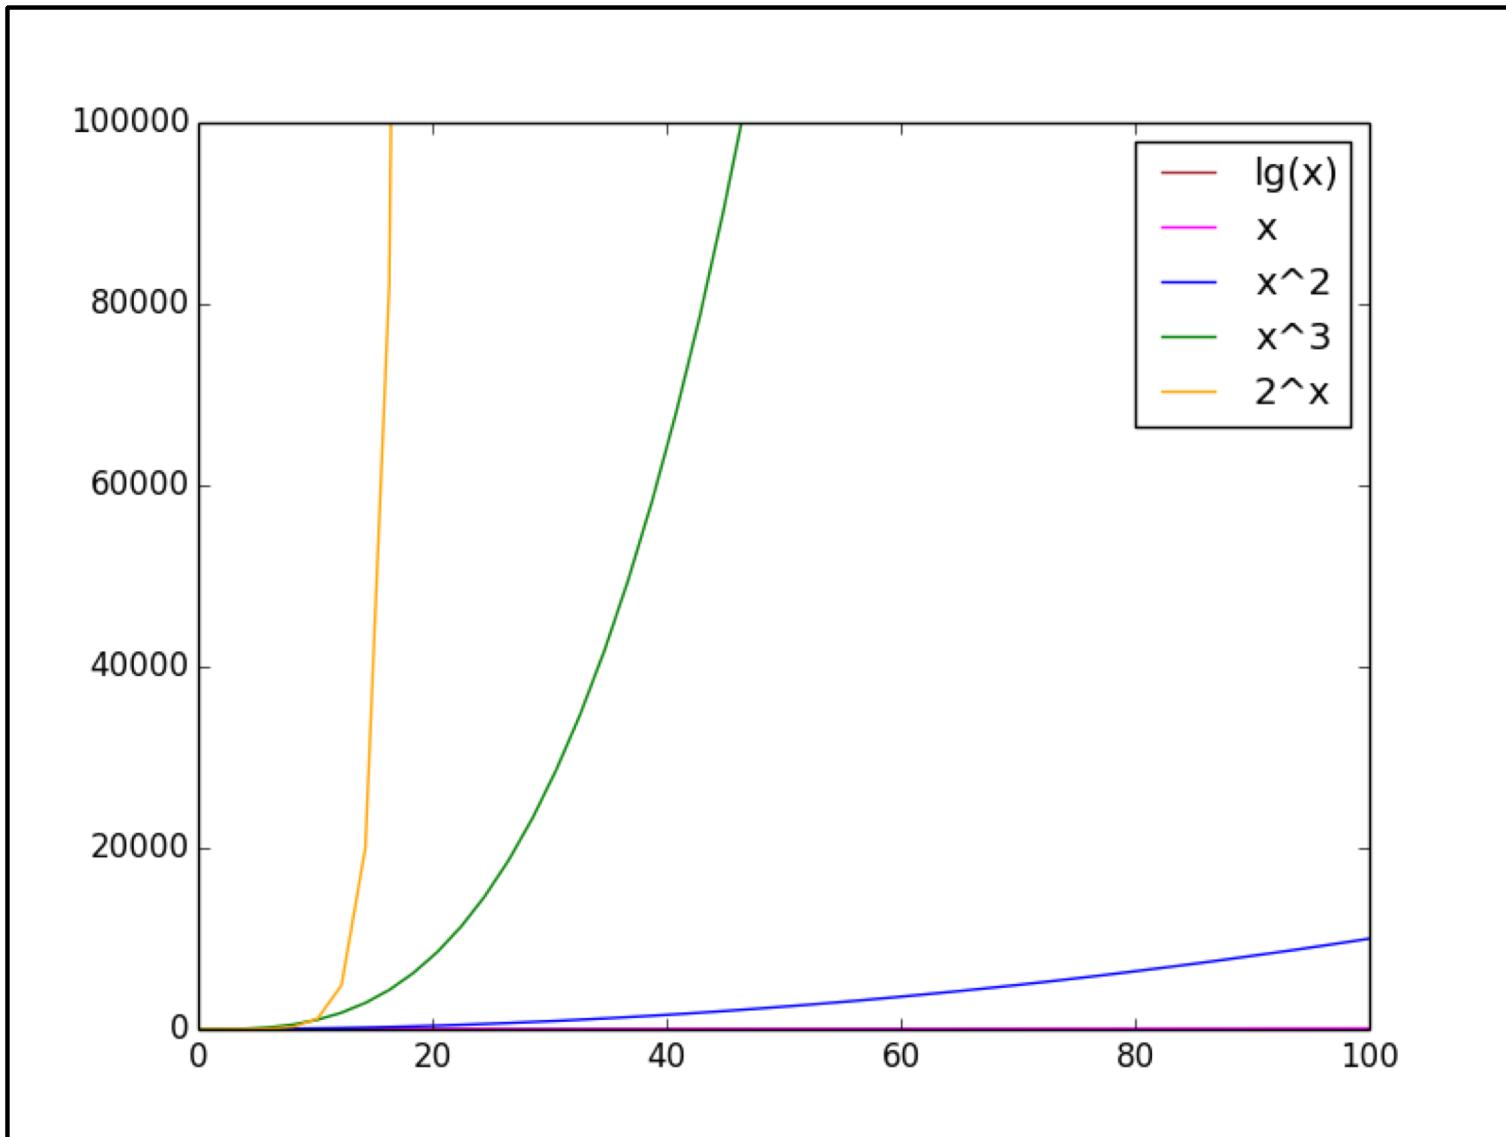
# Tilde approximations and Order of Growth

- ThreeSum order of growth:
  - Upper bound:  $O(n^3)$
  - Lower bound:  $\Omega(n^3)$
  - And hence:  $\Theta(n^3)$
- Tilde approximations?
  - Introduced in section 1.4 of the text
  - In this case:  $\sim n^3/6$

# Common orders of growth

- Constant - 1
- Logarithmic -  $\log n$
- Linear -  $n$
- Linearithmic -  $n \log n$
- Quadratic -  $n^2$
- Cubic -  $n^3$
- Exponential -  $2^n$
- Factorial -  $n!$

# Graphical orders of growth



How can we ignore lower order terms and multiplicative constants???

- Remember, this is asymptotic analysis

$f(n)$	$n =$	10	100	1,000	10,000
$n^3/6 - n^2/2 + n/3$		120	161,700	166,167,000	166,616,670,000
$n^3/6$		167	166,667	166,666,667	166,666,666,667
$n^3$		1,000	1,000,000	1,000,000,000	1,000,000,000,000

# Quick algorithm analysis

- Ignore multiplicative constants and lower terms
- Use standard measures for comparison

# Easy to get Theta for ThreeSum

- Why do we need to bother with Big O and Big Omega?

# Further thoughts on ThreeSum

- Is there a better way to solve the problem?
- What if we sorted the array first?
  - Pick two numbers, then binary search for the third one that will make a sum of zero
    - $a[i] = 10, a[j] = -7$ , binary search for -3
    - Still have two for loops, but we replace the third with a binary search
      - Runtime now?
      - What if the input data isn't sorted?
  - See ThreeSumFast.java

# Brief sorting review

- Given a list of  $n$  items, place the items in a given order
  - Ascending or descending
    - Numerical
    - Alphabetical
    - etc.

# Prerequisites

```
boolean less(Comparable v, Comparable w) {  
    return (v.compareTo(w) < 0);  
}
```

```
void exch(Object[] a, int i, int j) {  
    Object swap = a[i];  
    a[i] = a[j];  
    a[j] = swap;  
}
```

# Bubble sort

- Simply go through the array comparing pairs of items, swap them if they are out of order
  - Repeat until you make it through the array with 0 swaps

```
void bubbleSort(Comparable[] a) {  
    boolean swapped;  
    do {  
        swapped = false;  
        for(int j = 1; j < a.length; j++) {  
            if (less(a[j], a[j-1]))  
                { exch(a, j-1, j); swapped = true; }  
        }  
    } while(swapped);  
}
```

# Bubble sort example

**SWAPPED!**

1	3	4	5	10
---	---	---	---	----

# “Improved” bubble sort

```
void bubbleSort(Comparable[] a) {  
    boolean swapped;  
    int to_sort = a.length;  
    do {  
        swapped = false;  
        for(int j = 1; j < to_sort; j++) {  
            if (less(a[j], a[j-1]))  
                { exch(a, j-1, j); swapped = true; }  
        }  
        to_sort--;  
    } while(swapped);  
}
```

# How bad is it?

- Runtime:
  - $O(n^2)$

"[A]lthough the techniques used in the calculations [to analyze the bubble sort] are instructive, the results are disappointing since they tell us that the bubble sort isn't really very good at all."

Donald Knuth

*The Art of Computer Programming*



# Hybrid merge sort

- Read in amount of data that will fit in memory
- Sort it in place
  - i.e., via quick sort
- Write sorted chunk of data to disk
- Repeat until all data is stored in sorted chunks
- Merge chunks together

# Topics for the term

- Searching
  - Hashing
- Compression
- Large Integer Math
  - Cryptography
- Graph Algorithms
  - Heaps and Priority Queues
- P vs NP
  - Heuristic Approximation
- Dynamic Programming

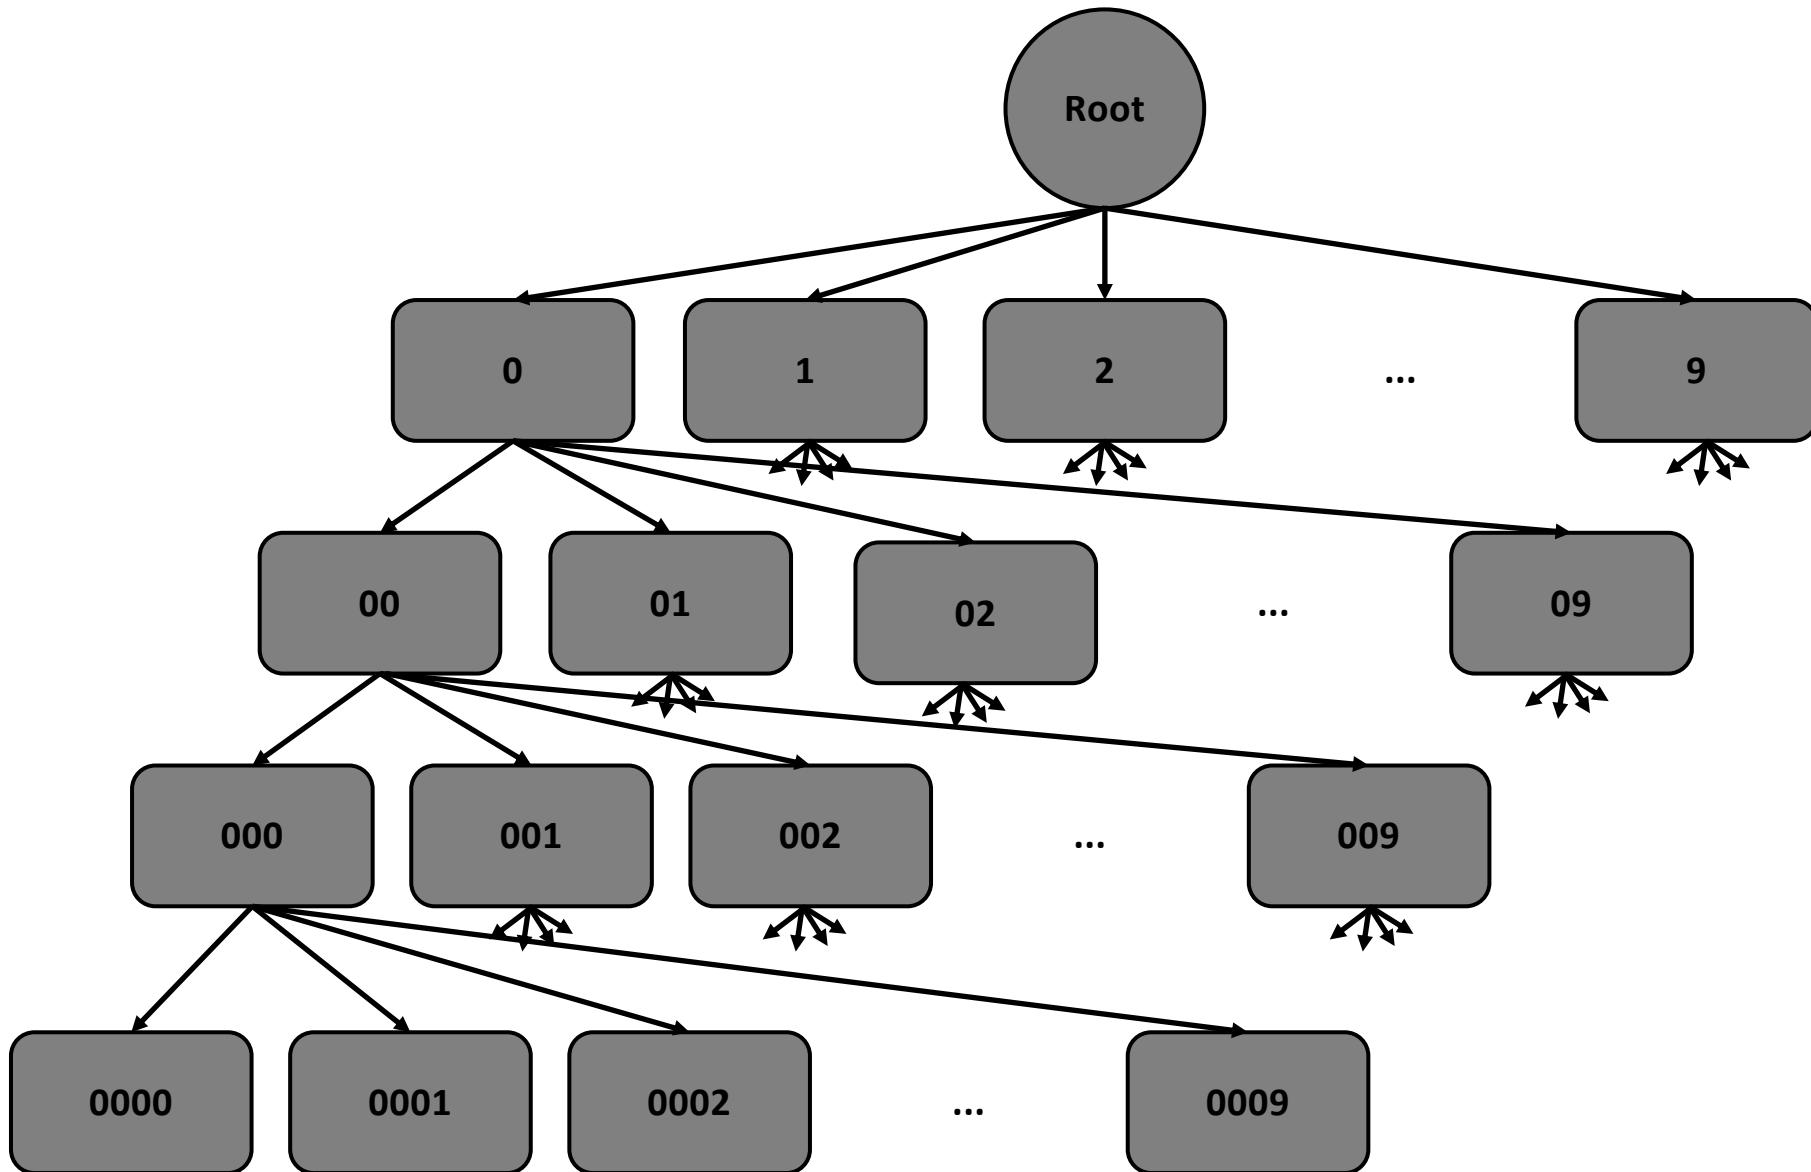
# Brute-force (or exhaustive) search

- Find the solution to a problem by considering all potential solutions and selecting the correct one
- Run-time is bounded by the number of potential solutions
  - $n^3$  potential solutions means cubic run-time
  - $2^n$  potential solutions means exponential run-time

# Password cracking

- Brute force password attacks depend on the length of the password, hence the insecurity of short passwords
- We can view the series of guesses we make as a tree
  - Each path from root to leaf is an attempted solution

# PIN cracking example

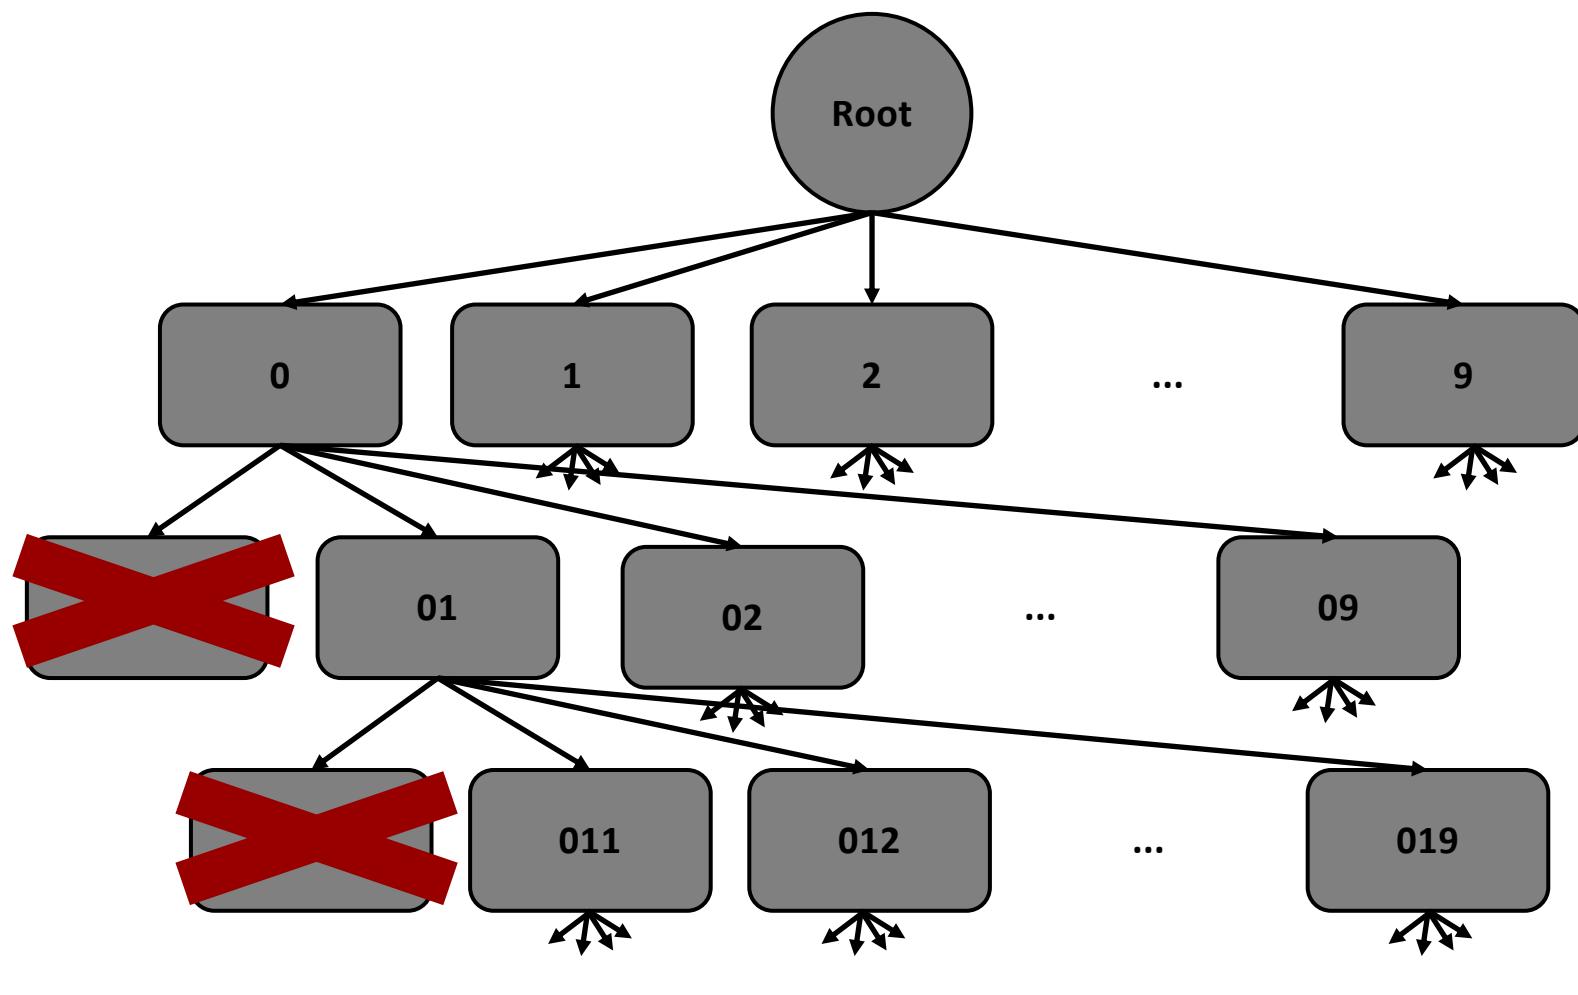


# Search space size

- This tree will enumerate  $10^n$  different PINs
  - $n$  is the length of the PIN
  - So for our case  $10^4 = 10,000$  different PINs
- Note that this is (for a computer) tiny
  - What would be a long password for a computer?
  - Say 128 bits long
    - $2^n$  different passwords
    - $2^{128} = 340282366920938463463374607431768211456$ 
      - Assuming a supercomputer can check 33860000000000 passwords per second...
      - And we'll on average find the correct password after guessing half the possibilities...
      - We should be able to crack a 128 bit password on our supercomputer in  $1.59 \times 10^{17}$  years using brute force

# Back to our PIN cracking example

- What if we have background knowledge that the PIN we're trying to crack doesn't have more than one 0?



# Pruning!

- Removes entire subtrees of our search space exploration
- When we can use it, it makes our algorithm practical for much larger values of n
- Does not, however, affect the asymptotic performance of an algorithm
  - Still exponential time requirement for our PIN example

# How to enumerate all these possibilities?

- For the PIN example a whole bunch of for loops would do
- In general, exhaustive search trees can be easily traversed via recursion and *backtracking*
  - Recurse until its apparent no solution can be achieved along the current path
  - Undo the path to the point that you can start to move forward again

# 8 queens problem

- Place 8 queens on a chessboard such that no queen can take another
  - Queens can move horizontally, vertically, and diagonally
  - How many ways can you places 8 pieces on a chess board?
    - $64 \text{ C } 8$
    - $= 64! / (8! * (64-8)!)$
    - $= 4,426,165,368$
    - Meaning 35,409,322,944 total queen placements
  - Do we really need to look through all of these options?

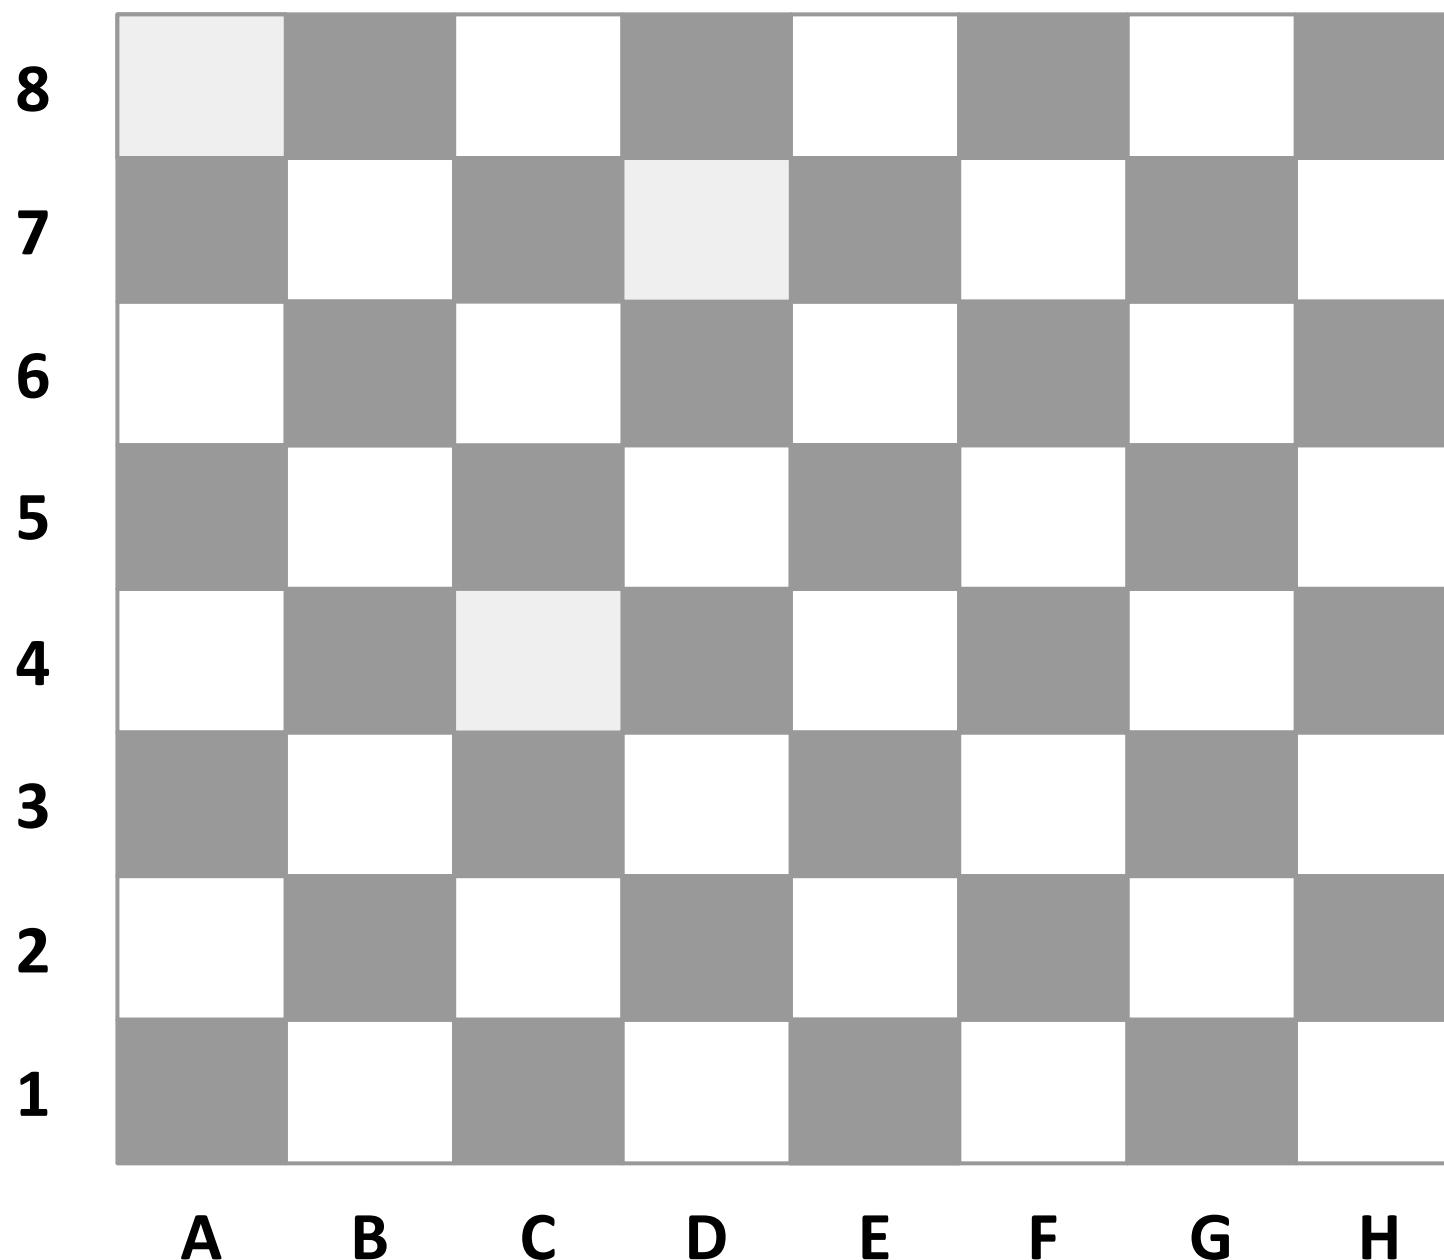
# 8 queens problem

- Solutions only have one queen per column
  - Still  $8^8 = 16,777,216$  possible combinations
- Solutions only have one queen per row
  - Combining these two observations, only  $8! = 40,320$
  - Looking quite feasible!
- Finally, prune subtrees with queens on the same diagonal

# Queens.java

- Basic idea:
  - Recurse over columns of the board
  - Each recursive call iterates through the rows of the board
  - Check rows/diagonals
    - Are they currently safe?
      - Place a queen in the current row/col
      - If you are at the end of the board, you've found a solution!
      - Otherwise, try recursive call for the next column
    - If they are not currently safe
      - Continue to the next row in the current recursive call

# 8 Queens



# Another problem: Boggle

- Words at least 3 adjacent letters long must be assembled from a 4x4 grid
- Adjacent letters are horizontally, vertically, or diagonally neighboring
- Any cube in the grid can only be used once per word



# Recurising through Boggle letters

- Have 8 different options from each cube
  - From  $B[i][j]$ :
    - $B[i-1][j-1]$
    - $B[i-1][j]$
    - $B[i-1][j+1]$
    - $B[i][j-1]$
    - $B[i][j+1]$
    - $B[i+1][j-1]$
    - $B[i+1][j]$
    - $B[i+1][j+1]$
  - Naively, the runtime here would be 16!
    - $= 20,922,789,888,000$

# Where do we prune?

# Implementation concerns with Boggle

- Constructing the words over the course of recursion will mean building up and tearing down strings
  - Moving forward adds a new character to the current word string
  - Backtracking removes the most recent character
  - Basically pushing/popping to/from a string stack
- Push/Pop stack operations are generally  $\Theta(1)$ 
  - Unless you need to resize, but that cost can be amortized
- Java Strings, however, are *immutable*
  - `s = new String("Here is a basic string");`
  - `s = s + " this operation allocates and initializes all over again";`
  - Becomes essentially a  $\Theta(n)$  operation
    - Where n is the length of the string

# StringBuilder to the rescue

- `append()` and `deleteCharAt()` can be used to push and pop
  - Back to  $\Theta(1)$ !
  - Still need to account for resizing, though...
- `StringBuffer` can also be used for this purpose
  - Differences?