



CS/COE 1501
Algorithm Implementation

4-Algorithms for Big Integers

Fall 2018

Sherif Khattab

ksm73@pitt.edu

6307 Sennott Square

(Slides are adapted from Dr. Ramirez's and Dr. Farnan's CS1501 slides.)

Integer multiplication

- Say we have 5 baskets with 8 apples in each
 - How do we determine how many apples we have?
 - Count them all?
 - That would take a while...
 - Since we know we have 8 in each basket, and 5 baskets, lets simply add $8 + 8 + 8 + 8 + 8$
 - $= 40!$
 - This is essentially multiplication!
 - $8 * 5 = 8 + 8 + 8 + 8 + 8$

What about bigger numbers?

- Like $1284 * 1583$, I mean!
 - That would take way longer than counting the 40 apples!
- Let's think of it like this:
 - $1284 * 1583 = 1284*3 + 1284*80 + 1284*500 + 1284*1000$

$$\begin{array}{r} 1284 \\ \times \quad 1583 \\ \hline 3852 \\ + \quad 102720 \\ + \quad 642000 \\ + \quad 1284000 \\ \hline = \quad 2032572 \end{array}$$

OK, I'm guessing we all knew that...

- ... and learned it quite some time ago ...
- So why bring it up now? What is there to cover about multiplication
- What is the runtime of this multiplication algorithm?
 - For 2 n-digit numbers:
 - n^2

Yeah, but the processor has a MUL instruction

- Assuming x86
- Given two 32-bit integers, MUL will produce a 64-bit integer in a few cycles
- What about when we need to multiply large ints?
 - VERY large ints?
 - RSA keys should be 2048 bits
 - Back to grade school...

Gradeschool algorithm on binary numbers

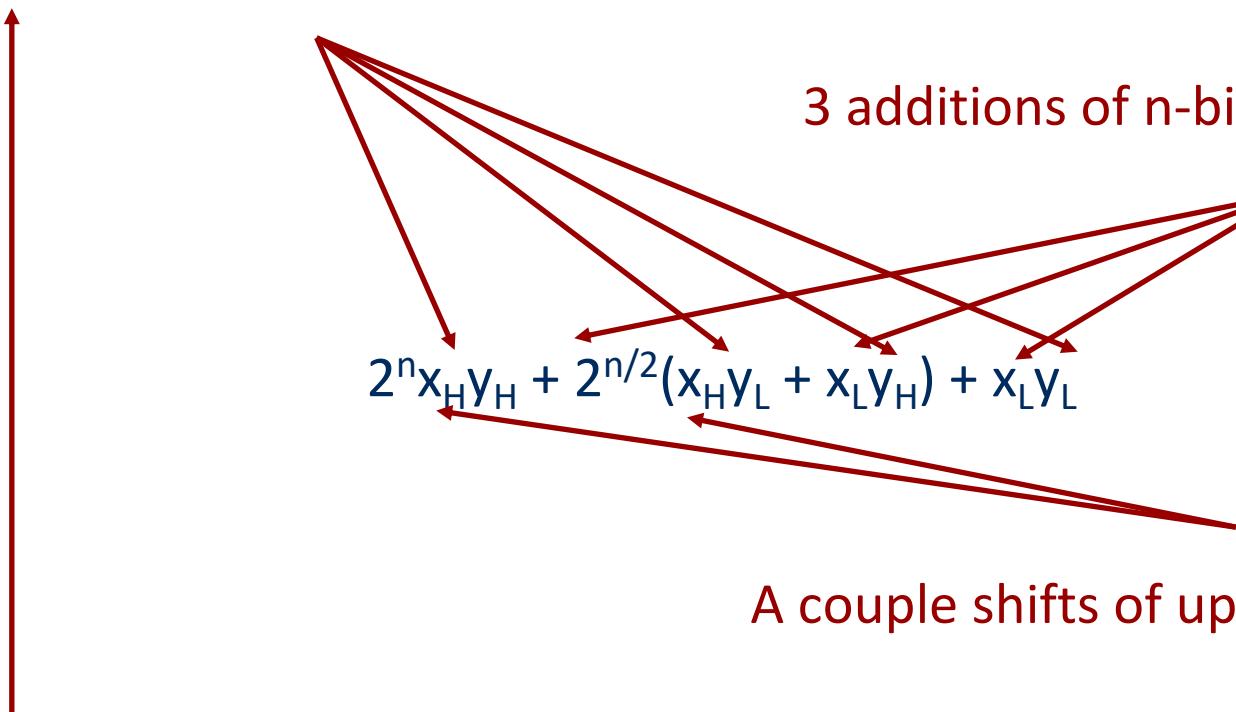
$$\begin{array}{r} 10100000100 \\ \times \quad 11000101111 \\ \hline 10100000100 \\ 101000001000 \\ 1010000010000 \\ 10100000100000 \\ 00000000000000 \\ 101000001000000 \\ 000000000000000 \\ 000000000000000 \\ 000000000000000 \\ 10100000100000000 \\ 10100000100000000 \\ \hline 11111000001110111100 \end{array}$$

How can we improve our runtime?

- Let's try to divide and conquer:
 - Break our n-bit integers in half:
 - $x = 1001011011001000$, $n = 16$
 - Let the high-order bits be $x_H = 10010110$
 - Let the low-order bits be $x_L = 11001000$
 - $x = 2^{n/2}x_H + x_L$
 - Do the same for y
 - $x * y = (2^{n/2}x_H + x_L) * (2^{n/2}y_H + y_L)$
 - $x * y = 2^n x_H y_H + 2^{n/2}(x_H y_L + x_L y_H) + x_L y_L$

So what does this mean?

4 multiplications of $n/2$ bit integers



Actually 16 multiplications of $n/4$ bit integers

(plus additions/shifts)

Actually 64 multiplications of $n/8$ bit integers

(plus additions/shifts)

...

So what's the runtime???

- Recursion really complicates our analysis...
- We'll use a *recurrence relation* to analyze the recursive runtime
 - Goal is to determine:
 - How much work is done in the current recursive call?
 - How much work is passed on to future recursive calls?
 - All in terms of input size

Recurrence relation for divide and conquer multiplication

- Assuming we cut integers exactly in half at each call
 - I.e., input bit lengths are a power of 2
- Work in the current call:
 - Shifts and additions are $\Theta(n)$
- Work left to future calls:
 - 4 more multiplications on half of the input size
- $T(n) = 4T(n/2) + \Theta(n)$

Sooooo... what's the runtime?

- Need to solve the recurrence relation
 - Remove the recursive component and express it purely in terms of n
 - A “cookbook” approach to solving recurrence relations:
 - The master theorem

The master theorem

- Usable on recurrence relations of the following form:

$$T(n) = aT(n/b) + f(n)$$

- Where:
 - a is a constant ≥ 1
 - b is a constant > 1
 - and $f(n)$ is an asymptotically positive function

Applying the master theorem

$$T(n) = aT(n/b) + f(n)$$

- If $f(n)$ is $O(n^{\log_b(a) - \varepsilon})$:
 - $T(n)$ is $\Theta(n^{\log_b(a)})$
- If $f(n)$ is $\Theta(n^{\log_b(a)})$
 - $T(n)$ is $\Theta(n^{\log_b(a)} \lg n)$
- If $f(n)$ is $\Omega(n^{\log_b(a) + \varepsilon})$ and $(a * f(n/b)) \leq c * f(n)$ for some $c < 1$:
 - $T(n)$ is $\Theta(f(n))$

Mergesort master theorem analysis

Recurrence relation for mergesort?

$$T(n) = 2T(n/2) + \Theta(n)$$

- $a = 2$
- $b = 2$
- $f(n) \text{ is } \Theta(n)$
- So...
 - $n^{\log_b(a)} = \dots$
 - $n^{\lg 2} = n$

- If $f(n)$ is $O(n^{\log_b(a) - \varepsilon})$:
 - $T(n) \text{ is } \Theta(n^{\log_b(a)})$
- If $f(n)$ is $\Theta(n^{\log_b(a)})$
 - $T(n) \text{ is } \Theta(n^{\log_b(a)} \lg n)$
- If $f(n)$ is $\Omega(n^{\log_b(a) + \varepsilon})$
and $(a * f(n/b)) \leq c * f(n)$ for some $c < 1$:
 - $T(n) \text{ is } \Theta(f(n))$

- Being $\Theta(n)$ means $f(n) \text{ is } \Theta(n^{\log_b(a)})$
- $T(n) = \Theta(n^{\log_b(a)} \lg n) = \Theta(n^{\lg 2} \lg n) = \Theta(n \lg n)$

For our divide and conquer multiplication approach

$$T(n) = 4T(n/2) + \Theta(n)$$

- $a = 4$
- $b = 2$
- $f(n)$ is $\Theta(n)$
- So...
 - $n^{\log_b a} = \dots$
 - $n^{\lg 4} = n^2$

- If $f(n)$ is $O(n^{\log_b a - \varepsilon})$:
 - $T(n)$ is $\Theta(n^{\log_b a})$
- If $f(n)$ is $\Theta(n^{\log_b a})$
 - $T(n)$ is $\Theta(n^{\log_b a} \lg n)$
- If $f(n)$ is $\Omega(n^{\log_b a + \varepsilon})$ and $(a * f(n/b)) \leq c * f(n)$ for some $c < 1$:
 - $T(n)$ is $\Theta(f(n))$

- Being $\Theta(n)$ means $f(n)$ is polynomially smaller than n^2
- $T(n) = \Theta(n^{\log_b a}) = \Theta(n^{\lg 4}) = \Theta(n^2)$

@#\$%^&*

- Leaves us back where we started with the grade school algorithm...
 - Actually, the overhead of doing all of the dividing and conquering will make it slower than grade school

SO WHY EVEN BOTHER?

- Let's look for a smarter way to divide and conquer
- Look at the recurrence relation again to see where we can improve our runtime:

$$T(n) = 4T(n/2) + \Theta(n)$$

Can we reduce the number of subproblems?

Can we reduce the amount of work done by the current call?

Can we reduce the subproblem size?

Karatsuba's algorithm

- By reducing the number of recursive calls (subproblems), we can improve the runtime
- $x * y = 2^n x_H y_H + 2^{n/2} (x_H y_L + x_L y_H) + x_L y_L$

M1 M2 M3 M4

- We don't actually need to do both M2 and M3
 - We just need the sum of M2 and M3
 - If we can find this sum using only 1 multiplication, we decrease the number of recursive calls and hence improve our runtime

Karatsuba craziness

- $M1 = x_h y_h; M2 = x_h y_l; M3 = x_l y_h; M4 = x_l y_l;$
- The sum of all of them can be expressed as a single mult:
 - $M1 + M2 + M3 + M4$
 - $= x_h y_h + x_h y_l + x_l y_h + x_l y_l$
 - $= (x_h + x_l) * (y_h + y_l)$
- Lets call this single multiplication $M5$:
 - $M5 = (x_h + x_l) * (y_h + y_l) = M1 + M2 + M3 + M4$
- Hence, $M5 - M1 - M4 = M2 + M3$
- So: $x * y = 2^n M1 + 2^{n/2}(M5 - M1 - M4) + M4$
 - Only 3 multiplications required!
 - At the cost of 2 more additions, and 2 subtractions

Karatsuba runtime

- To get M_5 , we have to multiply (at most) $n/2 + 1$ bit ints
 - Asymptotically the same as our other recursive calls
- Requires extra additions and subtractions...
 - But these are all $\Theta(n)$
- So, the recurrence relation for Karatsuba's algorithm is:
 - $T(n) = 3T(n/2) + \Theta(n)$
 - Which solves to be $\Theta(n^{\lg 3})$
 - Asymptotic improvement over grade school algorithm!
 - For large n , this will translate into practical improvement

Large integer multiplication in practice

- Can use a hybrid algorithm of grade school for large operands,
Karatsuba's algorithm for VERY large operands
 - Why are we still bothering with grade school at all?

Is this the best we can do?

- The Schönhage–Strassen algorithm
 - Uses Fast Fourier transforms to achieve better asymptotic runtime
 - $O(n \log n \log \log n)$
 - Fastest asymptotic runtime known from 1971–2007
 - Required n to be astronomical to achieve practical improvements to runtime
 - Numbers beyond $2^{2^{15}}$ to $2^{2^{17}}$
- Fürer was able to achieve even better asymptotic runtime in 2007
 - $n \log n 2^{O(\log^* n)}$
 - No practical difference for realistic values of n

Exponentiation

- x^y
- Can easily compute with a simple algorithm:

```
ans = 1
i = y
while i > 0:
    ans = ans * x
    i--
```

- Runtime?

Just like with multiplication, let's consider large integers...

- Runtime = # of iterations * cost to multiply
- Cost to multiply was covered in the last lecture
- So how many iterations?
 - Single loop from 1 to y , so linear, right?
 - What is the size of our input?
 - n
 - The *bitlength* of y ...
 - So, linear in the *value* of y ...
 - But, increasing n by 1 doubles the number of iterations
 - $\Theta(2^n)$
 - Exponential in the *bitlength* of y

This is RIDICULOUSLY BAD

- Assuming 512 bit operands, 2^{512} :
 - 1340780792994259709957402499820584612747936582059239337772356
1443721764030073546976801874298166903427690031858186486050853
753882811946569946433649006084096
- Assuming we can do mults in 1 cycle...
 - Which we *can't* as we learned last lecture
- And further that these operations are completely parallelizable
- 16 4GHz cores = 64,000,000,000 cycles/second
 - $(2^{512} / 64000000000) / (3600 * 24 * 365) =$
 - $6.64 * 10^{135}$ years to compute

This is way too long to do exponentiations!

- So how do we do better?
- Let's try divide and conquer!
- $x^y = (x^{(y/2)})^2$
 - When y is even, $(x^{(y/2)})^2 * x$ when y is odd
- Analyzing a recursive approach:
 - Base case?
 - When y is 1, x^y is x ; when y is 0, x^y is 1
 - Runtime?

Building another recurrence relation

- $x^y = (x^{(y/2)})^2 = x^{(y/2)} * x^{(y/2)}$
 - Similarly, $(x^{(y/2)})^2 * x = x^{(y/2)} * x^{(y/2)} * x$
- So, our recurrence relation is:
 - $T(n) = T(n-1) + ?$
 - How much work is done per call?
 - 1 (or 2) multiplication(s)
 - Examined runtime of multiplication last lecture
 - But how big are the operands in this case?

Determining work done per call

- Base case returns x
 - n bits
- Base case results are multiplied: $x * x$
 - n bit operands
 - Result size?
 - $2n$
- These results are then multiplied: $x^2 * x^2$
 - $2n$ bit operands
 - Result size?
 - $4n$ bits
- ...
- $x^{(y/2)} * x^{(y/2)}$?
 - $(y / 2) * n$ bit operands = $2^{(n-1)} * n$ bit operands
 - Result size? $y * n$ bits = $2^n * n$ bits

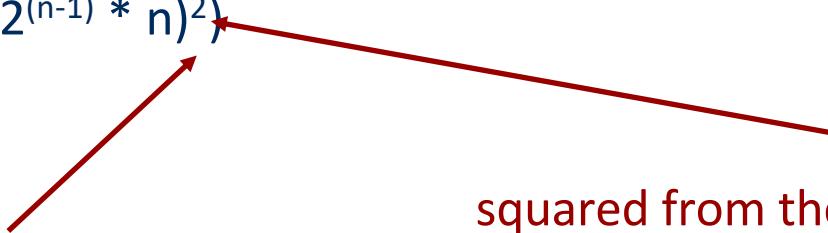
Multiplication input size increases throughout

- Our recurrence relation looks like:

- $T(n) = T(n-1) + \Theta((2^{(n-1)} * n)^2)$

multiplication input size

squared from the use of the
gradeschool algorithm



Runtime analysis

- Can we use the master theorem?
 - Nope, we don't have a $b > 1$
- OK, then...
 - How many times can y be divided by 2 until a base case?
 - $\lg(y)$
 - Further, we know the max value of y
 - Relative to n , that is:
 - 2^n
 - So, we have, at most $\lg(y) = \lg(2^n) = n$ recursions

But we need to do expensive mult in each call

- We need to do $\Theta((2^{(n-1)} * n)^2)$ work in just the root call!
 - Our runtime is dominated by multiplication time
 - Exponentiation quickly generates HUGE numbers
 - Time to multiply them quickly becomes impractical

Can we do better?

- We go “top-down” in the recursive approach
 - Start with y
 - Halve y until we reach the base case
 - Square base case result
 - Continue combining until we arrive at the solution
- What about a “bottom-up” approach?
 - Start with our base case
 - Operate on it until we reach a solution

A bottom-up approach

- To calculate x^y

```
ans = 1
foreach bit in y: ← From most to least significant
    ans = ans2
    if bit == 1:
        ans = ans * x
```

Bottom-up exponentiation example

- Consider x^y where y is 43 (computing x^{43})
- Iterate through the bits of y (43 in binary: 101011)
- $\text{ans} = 1$

$$\text{ans} = 1^2 \quad = 1$$

$$\text{ans} = 1 * x \quad = x$$

$$\text{ans} = x^2 \quad = x^2$$

$$\text{ans} = (x^2)^2 \quad = x^4$$

$$\text{ans} = x^4 * x \quad = x^5$$

$$\text{ans} = (x^5)^2 \quad = x^{10}$$

$$\text{ans} = (x^{10})^2 \quad = x^{20}$$

$$\text{ans} = x^{20} * x \quad = x^{21}$$

$$\text{ans} = (x^{21})^2 \quad = x^{42}$$

$$\text{ans} = x^{42} * x \quad = x^{43}$$

Does this solve our problem with mult times?

- Nope, still squaring ans everytime
 - We'll have to live with huge output sizes
- This does, however, save us recursive call overhead
 - Practical savings in runtime

Greatest Common Divisor

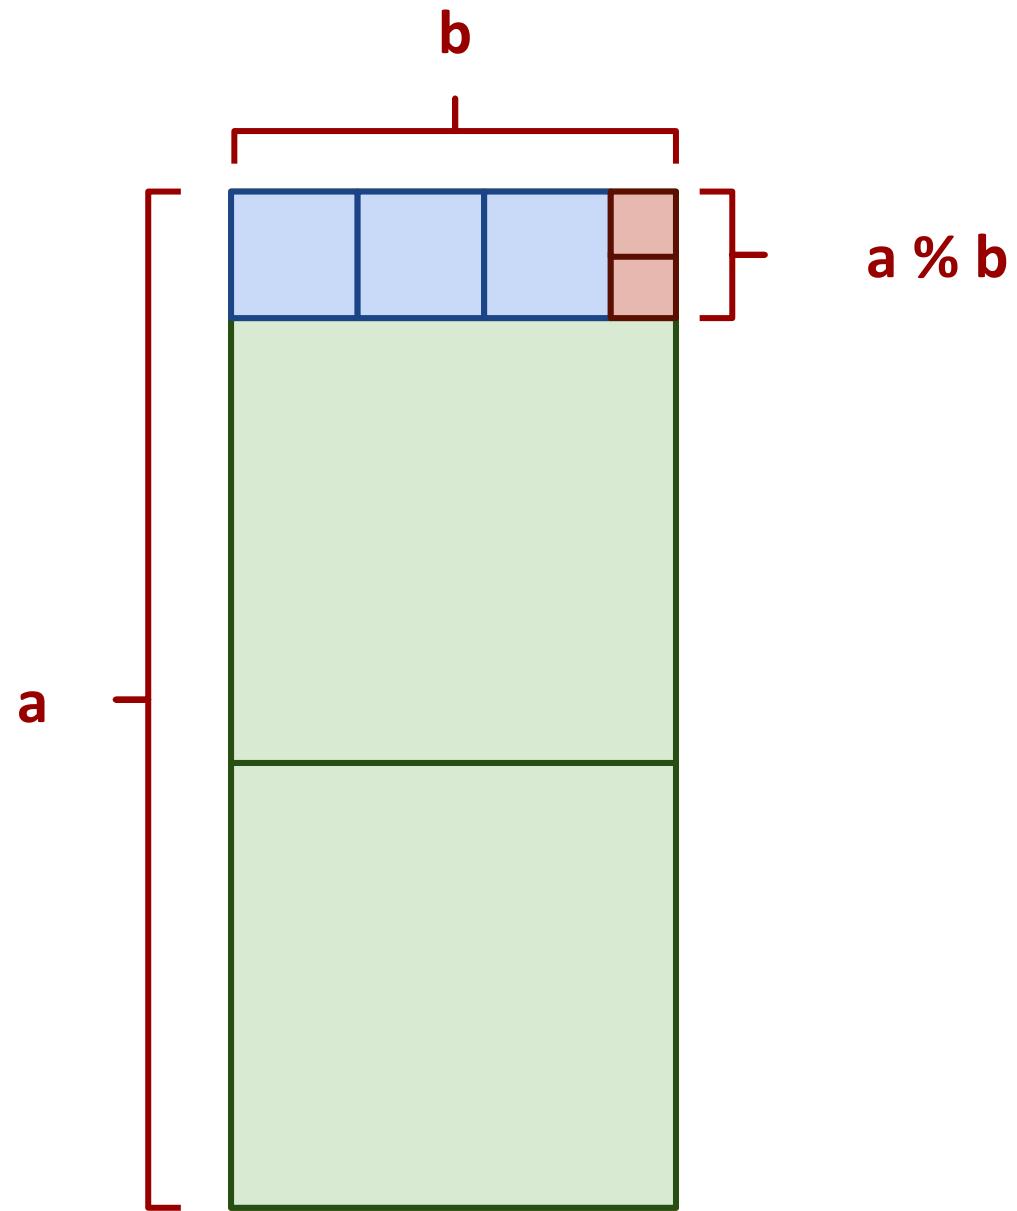
- $\text{GCD}(a, b)$
 - Largest int that evenly divides both a and b
- Easiest approach:
 - BRUTE FORCE

```
i = min(a, b)
while(a % i != 0 || b % i != 0):
    i--
```

- Runtime?
 - $\Theta(\min(a, b))$
 - Linear!
 - In *value* of $\min(a, b)$...
 - Exponential in n
 - Assuming a, b are n-bit integers

Euclid's algorithm

- $\text{GCD}(a, b) = \text{GCD}(b, a \% b)$



Euclidean example 1

- $\text{GCD}(30, 24)$
 - $= \text{GCD}(24, 30 \% 24)$
- $= \text{GCD}(24, 6)$
 - $= \text{GCD}(6, 24 \% 6)$
- $= \text{GCD}(6, 0)...$
 - Base case! Overall GCD is 6

Euclidean example 2

- $= \text{GCD}(99, 78)$
 - $99 = 78 * 1 + 21$
- $= \text{GCD}(78, 21)$
 - $78 = 21 * 3 + 15$
- $= \text{GCD}(21, 15)$
 - $21 = 15 * 1 + 6$
- $= \text{GCD}(15, 6)$
 - $15 = 6 * 2 + 3$
- $= \text{GCD}(6, 3)$
 - $6 = 3 * 2 + 0$
- $= 3$

Analysis of Euclid's algorithm

- Runtime?
 - Tricky to analyze, has been shown to be linear in n
 - Where, again, n is the number of bits in the input

Extended Euclidean algorithm

- In addition to the GCD, the Extended Euclidean algorithm (XGCD) produces values x and y such that:
 - $\text{GCD}(a, b) = i = ax + by$
- Examples:
 - $\text{GCD}(30, 24) = 6 = 30 * 1 + 24 * -1$
 - $\text{GCD}(99, 78) = 3 = 99 * -11 + 78 * 14$
- Can be done in the same linear runtime!

Extended Euclidean example

- $= \text{GCD}(99, 78)$
 - $99 = 78 * 1 + 21$
 - $= \text{GCD}(78, 21)$
 - $78 = 21 * 3 + 15$
 - $= \text{GCD}(21, 15)$
 - $21 = 15 * 1 + 6$
 - $= \text{GCD}(15, 6)$
 - $15 = 6 * 2 + 3$
 - $= \text{GCD}(6, 3)$
 - $6 = 3 * 2 + 0$
 - $= 3$
-
- $3 = 15 - (2 * 6)$
 - $6 = 21 - 15$
 - $3 = 15 - (2 * (21 - 15))$
 - $= 15 - (2 * 21) + (2 * 15)$
 - $= (3 * 15) - (2 * 21)$
 - $15 = 78 - (3 * 21)$
 - $3 = (3 * (78 - (3 * 21))) - (2 * 21)$
 - $= (3 * 78) - (11 * 21)$
 - $21 = 99 - 78$
 - $3 = (3 * 78) - (11 * (99 - 78))$
 - $= (14 * 78) - (11 * 99)$
 - $= 99 * -11 + 78 * 14$

OK, but why?

- This and all of our large integer algorithms will be handy when we look at algorithms for implementing...

CRYPTOGRAPHY

Introduction to crypto

- Cryptography - enabling secure communication in the presence of third parties
 - Alice wants to send Bob a message without anyone else being able to read it



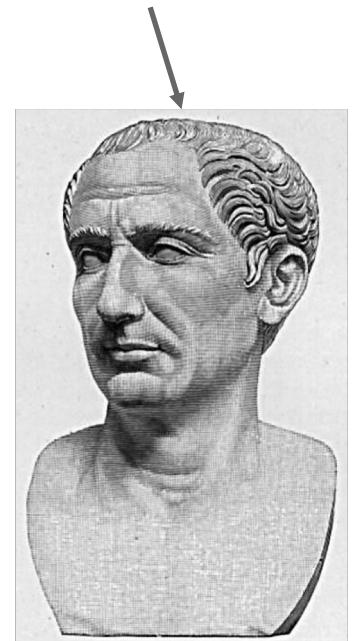
Enter the adversary

- Consider the adversary to be anyone that could try to eavesdrop on Alice and Bob communicating
 - People in the same coffee shop as Alice or Bob as they talk over WiFi
 - Admins operating the network between Alice and Bob
 - And mirroring their traffic to the NSA...
- Will have access to:
 - The *ciphertext*
 - The encrypted message
 - The encryption algorithm
 - At least Alice and Bob should assume the adversary does
- The key material is the only thing Bob knows that the adversary does not

Cryptography has been around for some time

- Early, classic encryption scheme:
 - Caesar cipher:
 - “Shift” the alphabet by a set amount
 - Use this shifted alphabet to send messages
 - The “key” is the amount the alphabet is shifted

Yes, that Caesar



Alphabet

ABCDEFGHIJKLMNOPQRSTUVWXYZ
XYZABCDEFGHIJKLM NOPQRSTUVWXYZ

Shift 3

By modern standards, incredibly easy to crack

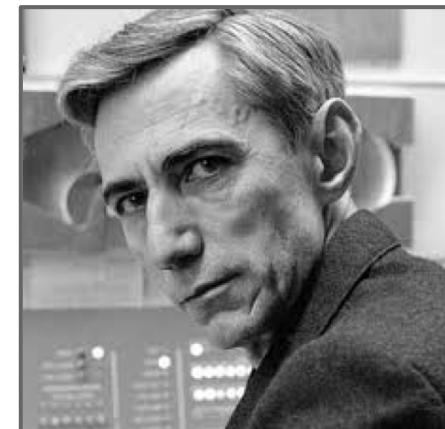
- BRUTE FORCE
 - Try every possible shift
 - 25 options for the English alphabet
 - 255 for ASCII
- OK, let's make it harder to brute force
 - Instead of using a shifted alphabet, let's use a random permutation of the alphabet
 - Key is now this permutation, not just a shift value
 - R size alphabet means $R!$ possible permutations!

By modern standards, incredibly easy to crack

- Just requires a bit more sophisticated of an algorithm
- Analyzing encrypted English for example
 - Sentences have a given structure
 - Character frequencies are skewed
 - Essentially playing Wheel of Fortune

So what is a good cipher?

- One-time pads
 - List of one-time use keys (called a *pad*) here
- To send a message:
 - Take an unused pad
 - Use modular addition to combine key with message
 - For binary data, XOR
 - Send to recipient
- Upon receiving a message:
 - Take the next pad
 - Use modular subtraction to combine key with message
 - For binary data, XOR
 - Read result
- Proven to provide perfect secrecy



One-time pad example

Encoding:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25
A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z

Message:

H E L L O
7 4 11 11 14

Pad: Q J C W T
16 9 2 22 19

$$\begin{array}{r} + \quad 16 \quad 9 \quad 2 \quad 22 \quad 19 \\ \hline \end{array} \quad (\text{mod } 26)$$

23 13 13 7 7

Encrypted Message:

X N N H H
23 13 13 7 7

$$\begin{array}{r} - \quad 16 \quad 9 \quad 2 \quad 22 \quad 19 \\ \hline \end{array} \quad (\text{mod } 26)$$

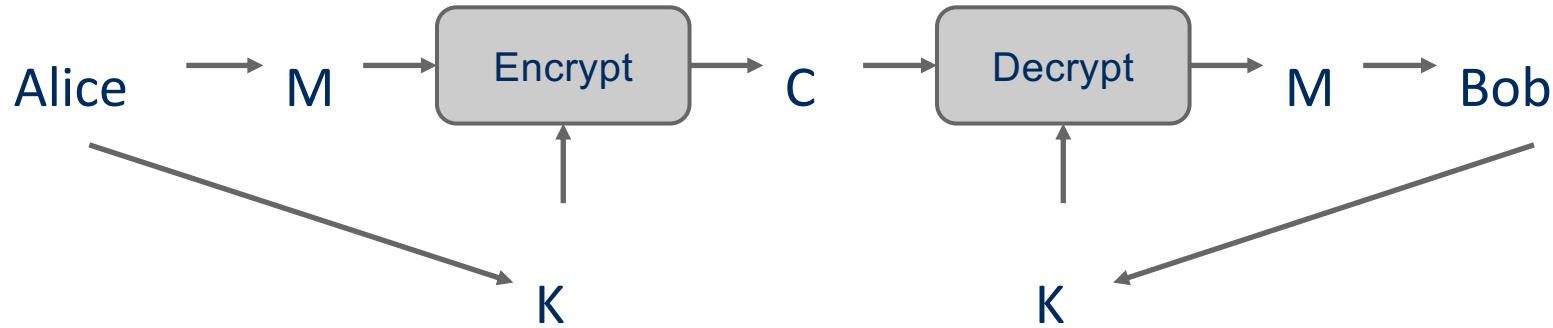
7 4 11 11 14

H E L L O

Difficulties with one-time pads

- Pads must be truly random
- Both sender and receiver must have a matched list of pads in the appropriate order
- Once you run out of pads, no more messages can be sent

Symmetric ciphers



- E.g., DES, AES, Blowfish
- Users share a single *key*
 - Numbers of a given bitlength (e.g., 128, 256)
 - Key is used to encrypt/decrypt many messages back and forth
- Encryptions/decryptions will be fast
 - Typically linear in the size the input
- Ciphertext should appear random
- Best way to recover plaintext should be a brute force attack on the encryption key
 - Which we have shown to be infeasible for 128bit AES keys

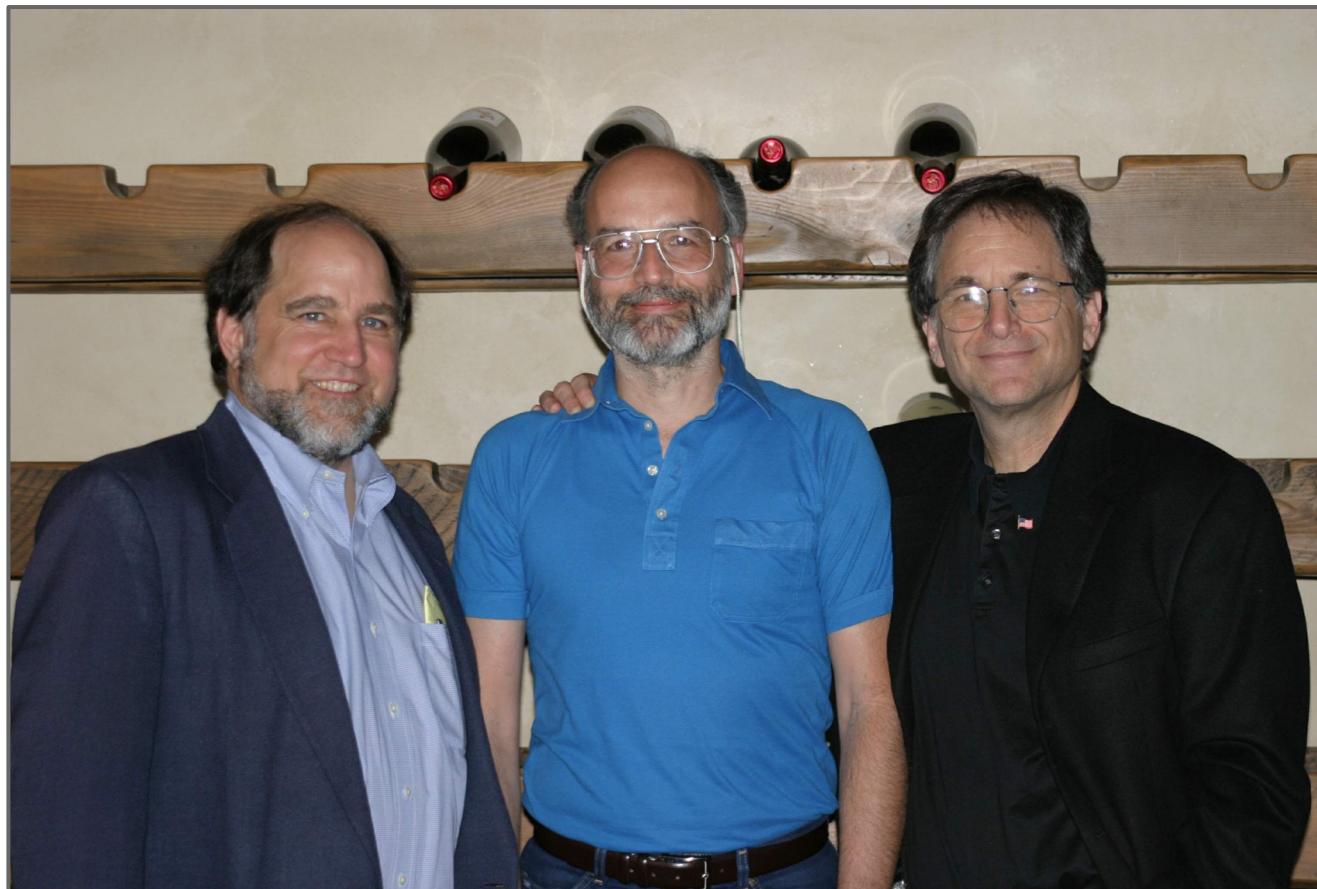
Problems with symmetric ciphers

- Alice and Bob have to both know the same key
 - How can you securely transmit the key from Alice to Bob?
- Further, if Alice also wants to communicate with Charlie, her and Charlie will need to know the same key, a different key from the key Alice shares with Bob
 - Alice and Danielle will also have to share a different key...
 - etc.

Enter public-key encryption

- Each user has their own pair of keys
 - A *public* key that can be revealed to anyone
 - A *private* key that only they should know
- How does this solve our problem?
 - Public key can simply be published/advertised
 - Posted repositories of public keys
 - Added to an email signature
 - Each user is responsible only for their own keypair
- Let's look at a public-key crypto scheme in detail...

RSA



RSA Cryptosystem in-depth

- What are RSA keypairs?
- How messages encrypted?
- How are messages decrypted?
- How are keys generated?
- Why is it secure?

RSA keypairs

- *Public* key is two numbers, which we will call **n** and **e**
- *Private* key is a single number we will call **d**
- The length of **n** in bits is the key length
 - I.e., 2048 bit RSA keys will have a 2048 bit **n** value
 - Note that "n" will be used to indicate the RSA public key component for our discussion of RSA...

Encryption

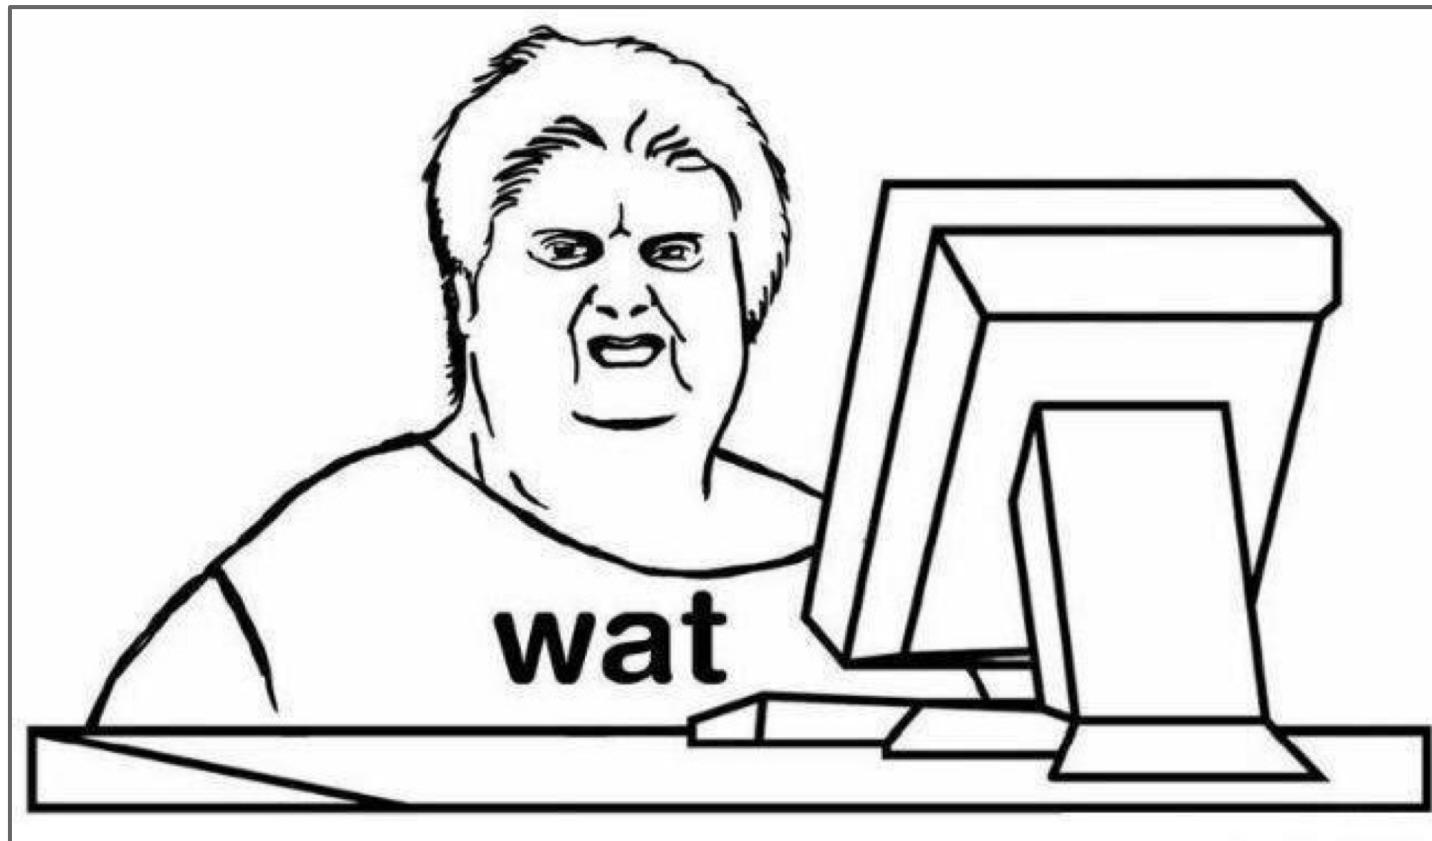
Say Alice wants to send a message to Bob

1. Looks up Bob's public key
2. Convert the message into an integer: m
3. Compute the ciphertext c as:
 - o $c = m^e \pmod{n}$
4. Send c to Bob

Decryption

Bob can simply:

1. Compute m as:
 - a. $m = c^d \pmod{n}$
2. Convert m into Alice's message



n , e , and d need to be carefully generated

1. Choose two prime numbers p and q
2. Compute $n = p * q$
3. Compute $\phi(n)$
 - o $\phi(n) = \phi(p) * \phi(q) = (p - 1) * (q - 1)$
4. Choose e such that
 - o $1 < e < \phi(n)$
 - o $\text{GCD}(e, \phi(n)) = 1$
 - I.e., e and $\phi(n)$ are co-prime
5. Determine d as $d = e^{-1} \bmod(\phi(n))$

What the ϕ ?

- Here, we mean ϕ to be Euler's totient
- $\phi(n)$ is a count of the integers $< n$ that are co-prime to n
 - I.e., how many k are there such that:
 - $1 \leq k \leq n$ AND $\text{GCD}(n, k) = 1$
- p and q are prime..
 - Hence, $\phi(p) = p - 1$ and $\phi(q) = q - 1$
- Further, ϕ is multiplicative
 - Since p and q are prime, they are co-prime, so
 - $\phi(p) * \phi(q) = \phi(p * q) = \phi(n)$
 - I won't detail the proof here...

OK, now what about multiplicative inverses mod $\phi(n)$?

- $d = e^{-1} \text{ mod}(\phi(n))$
- Means that $d = 1/e \text{ mod}(\phi(n))$
- Means that $e * d = 1 \pmod{\phi(n)}$
- Now, *this* can be equivalently stated as $e * d = z * \phi(n) + 1$
 - For some z
- Can further restate this as: $e * d - z * \phi(n) = 1$
- Or similarly: $1 = \phi(n) * (-z) + e * d$
- How can we solve this?
 - Hint: recall that we know $\text{GCD}(\phi(n), e) = 1$

Use extended Euclidean algorithm!

- $\text{GCD}(a, b) = i = ax + by$
- Let:
 - $a = \phi(n)$
 - $b = e$
 - $x = -z$
 - $y = d$
 - $i = 1$
- $\text{GCD}(\phi(n), e) = 1 = \phi(n) * (-z) + e * d$
- We can compute d in linear time!

RSA keypair example notes

- p and q must be prime
- $n = p * q$
- $\phi(n) = (p - 1) * (q - 1)$
- Choose e such that
 - $1 < e < \phi(n)$ and $\text{GCD}(e, \phi(n)) = 1$
- Solve $\text{XGCD}(\phi(n), e) = 1 = \phi(n) * (-z) + e * d$
- Compute the ciphertext c as:
 - $c = m^e \pmod{n}$
- Recover m as:
 - $m = c^d \pmod{n}$

OK, but how does $m^{ed} = m \bmod n$?

- Feel free to look up the proof using Fermat's little theorem
 - Knowing this proof is **NOT** required for the course
 - Knowing how to generate RSA keys and encrypt/decrypt **IS**
- For this course, we'll settle with our example showing that it *does* work

Why is RSA secure?

- 4 avenues of attack on the math of RSA were identified in the original paper:
 - Factoring n to find p and q
 - Determining $\phi(n)$ without factoring n
 - Determining d without factoring n or learning $\phi(n)$
 - Learning to take e^{th} roots modulo n

Factoring n

- To the best of our knowledge, this is *hard*
 - A 768 bit RSA key was factored one time using the best currently known algorithm
 - Took 1500 CPU years
 - 2 years of real time on hundreds of computers
 - Hence, large keys are pretty safe
 - 2048 bit keys are a pretty good bet for now

What about determining $\phi(n)$ without factoring n?

- Would allow us to easily compute d because $ed = 1 \bmod \phi(n)$
- Note:
 - $\phi(n) = n - p - q + 1$
 - $\phi(n) = n - (p + q) + 1$
 - $(p + q) = n + 1 - \phi(n)$
 - $(p + q) - (p - q) = 2q$
 - Now we just need $(p - q)...$
 - $(p - q)^2 = p^2 - 2pq + q^2$
 - $(p - q)^2 = p^2 + 2pq + q^2 - 4pq$
 - $(p - q)^2 = (p + q)^2 - 4pq$
 - $(p - q)^2 = (p + q)^2 - 4n$
 - $(p - q) = \sqrt{((p + q)^2 - 4n)}$
- If we can figure out $\phi(n)$ efficiently, we could factor n efficiently!

Determining d without factoring n or learning $\phi(n)$?

- If we know, d, we can get a multiple of $\phi(n)$
 - $ed = 1 \bmod \phi(n)$
 - $ed = k\phi(n) + 1$
 - For some k
 - $ed - 1 = k\phi(n)$
- It has been shown that n can be efficiently factored using any multiple of $\phi(n)$
 - Hence, this would provide another efficient solution to factoring!

Learning to take e^{th} roots modulo n

- Conjecture was made in 1978 that breaking RSA would yield an efficient factoring algorithm
 - To date, it has been not been proven or disproven

This all leads to the following conclusion

- Odds are that breaking RSA efficiently implies that factoring can be done efficiently.
- Since factoring is probably hard, RSA is probably safe to use.

Implementation concerns

- Encryption/decryption:
 - How can we perform efficient exponentiations?
- Key generation:
 - We can do multiplication, XGCD for large integers
 - What about finding large prime numbers?

Efficient exponentiation for RSA

Does this solve our problems?

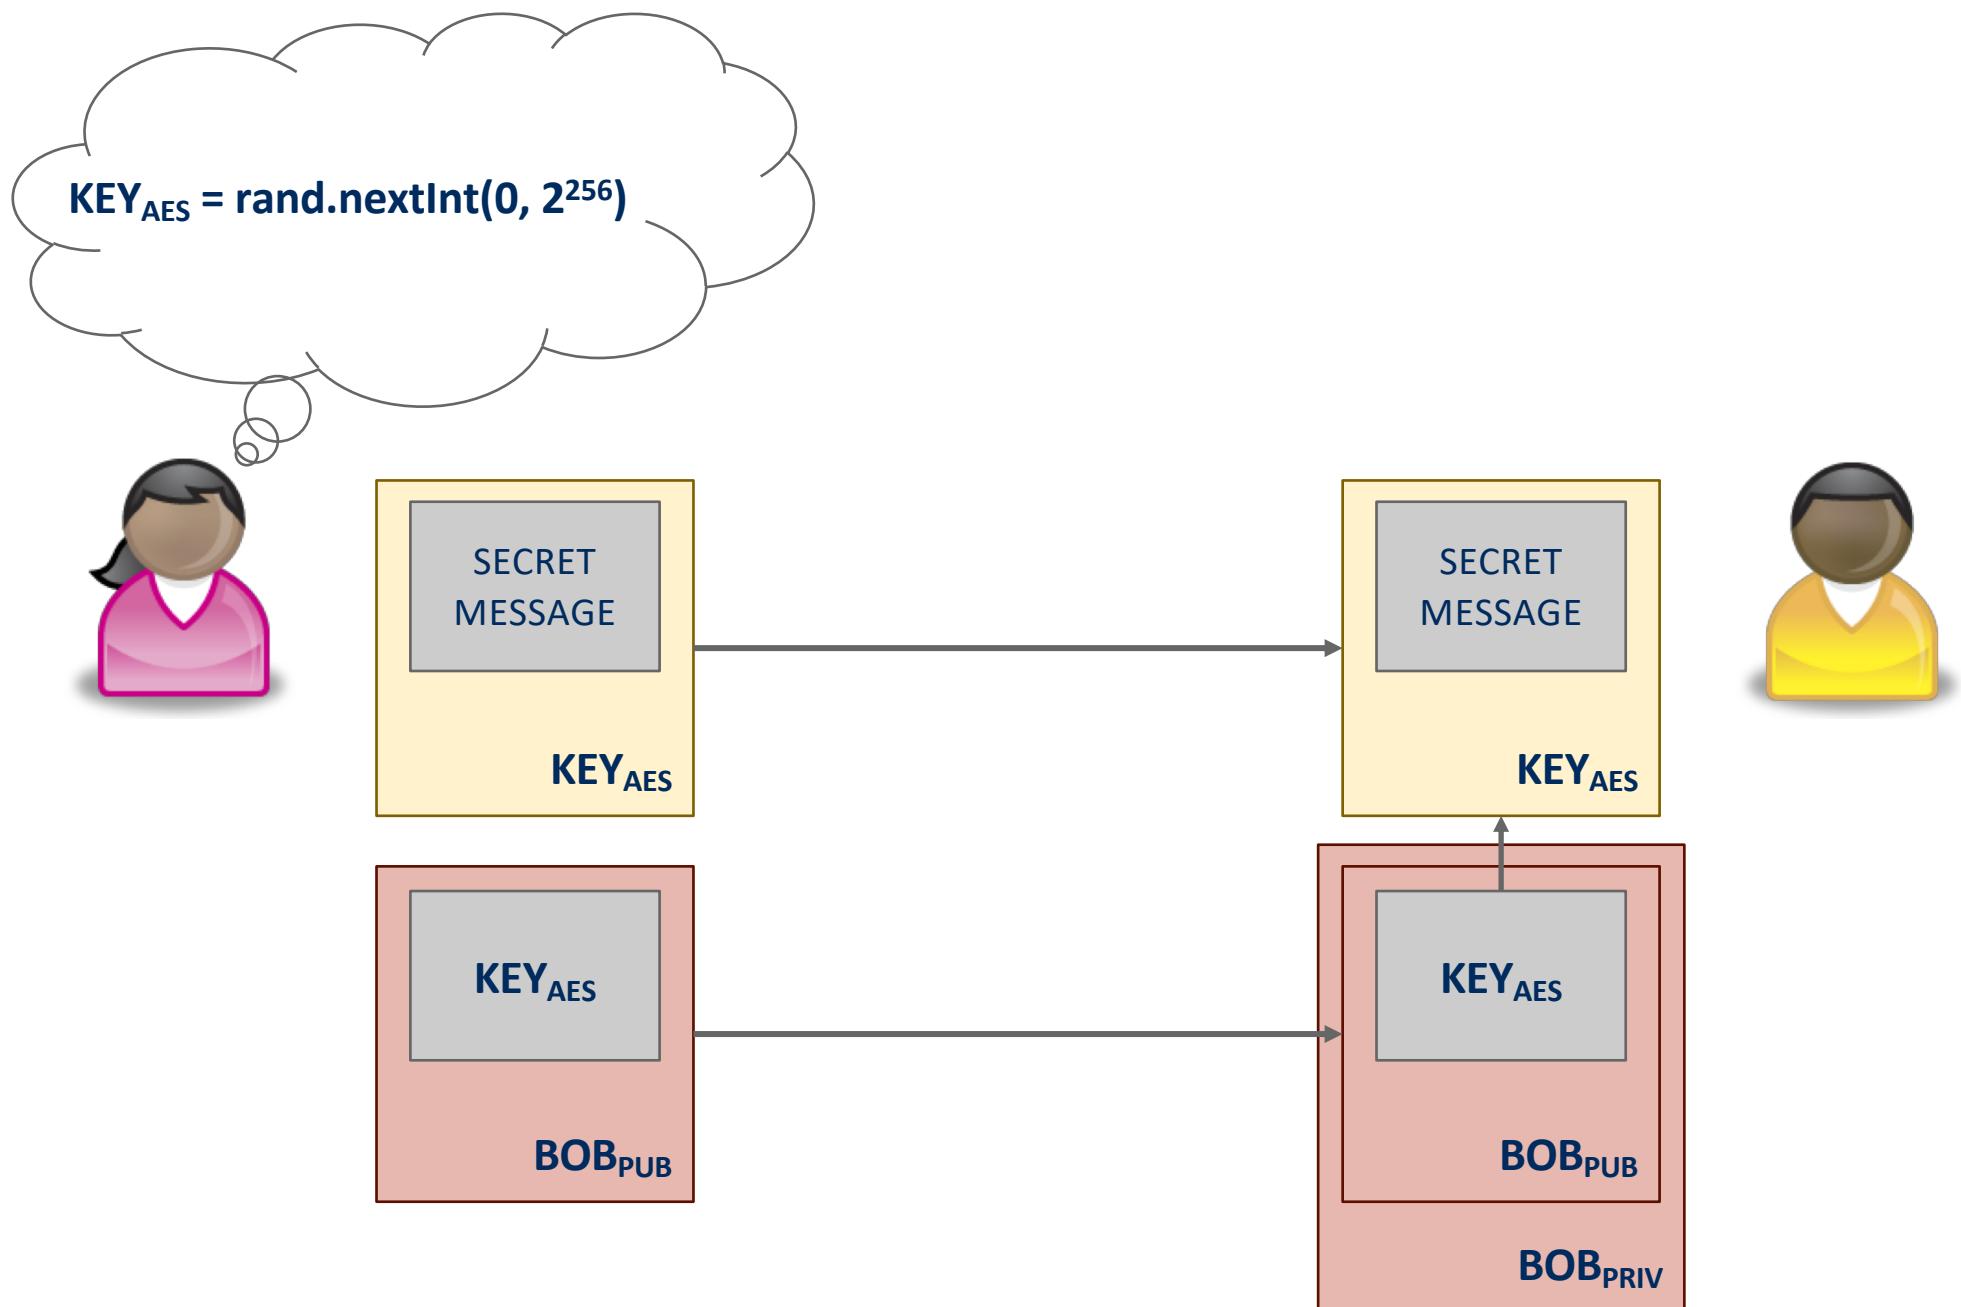
```
ans = 1
foreach bit in y:
    ans = (ans2 mod n)
    if bit == 1:
        ans = (ans * x
                mod n)x
```

- How can we improve runtime for RSA exponentiations?
 - Don't actually need x^y
 - Just need $(x^y \bmod n)$

Still slower (generally) than symmetric encryption

- If only we could have the speed of symmetric encryption without the key distribution woes!
 - What if we transmitted symmetric crypto keys with RSA?
 - RSA Envelopes!
- Going back to Alice and Bob
 - Alice generates a random AES key
 - Alice encrypts her message using AES with this key
 - Alice encrypts the key using Bob's RSA public key
 - Alice sends the encrypted message and encrypted key to Bob
 - Bob decrypts the AES key using his RSA private key
 - Bob decrypts the message using the AES key

RSA Envelope example



Prime testing option 1: BRUTE FORCE

- Try all possible factors of x
 - $1 \dots \text{sqrt}(x)$
 - aka $1 \dots \text{sqrt}(2^{\text{size}(x)})$
 - For a total of $2^{(\text{size}(x)/2)}$ factor checks
- A factor check should take about the same amount of time as multiplication
 - $\text{size}(x)^2$
- So our runtime is $\Theta(2^{(\text{size}(x)/2)} * \text{size}(x)^2)$

Option 2: A probabilistic approach

- Need a method test : $Z \times Z \rightarrow \{T, F\}$
 - If $\text{test}(x, a) = F$, x is composite based on the witness a
 - If $\text{test}(x, a) = T$, x is probably prime based on the witness a
 - To test a number x for primality:
 - Randomly choose a witness a
 - if $\text{test}(x, a) = F$, x is composite
 - if $\text{test}(x, a) = T$, loop
 - Possible implementations of $\text{test}(x, a)$:
 - Miller-Rabin, Fermat's, Solovay–Strassen
- often probability $\approx 1/2$
- k repetitions leads to probability that x is composite $\approx 1/2^k$

Another fun use of RSA...

- Notice that encrypting and decrypting are inverses
 - $m^{ed} = m^{de} \pmod{n}$
- We can “decrypt” the message first with a private key
- Then recover the message by “encrypting” with a public key
- Note that anyone can recover the message
 - However, they know the message must have come from the owner of the private key
 - Using RSA this way creates a *digital signature*

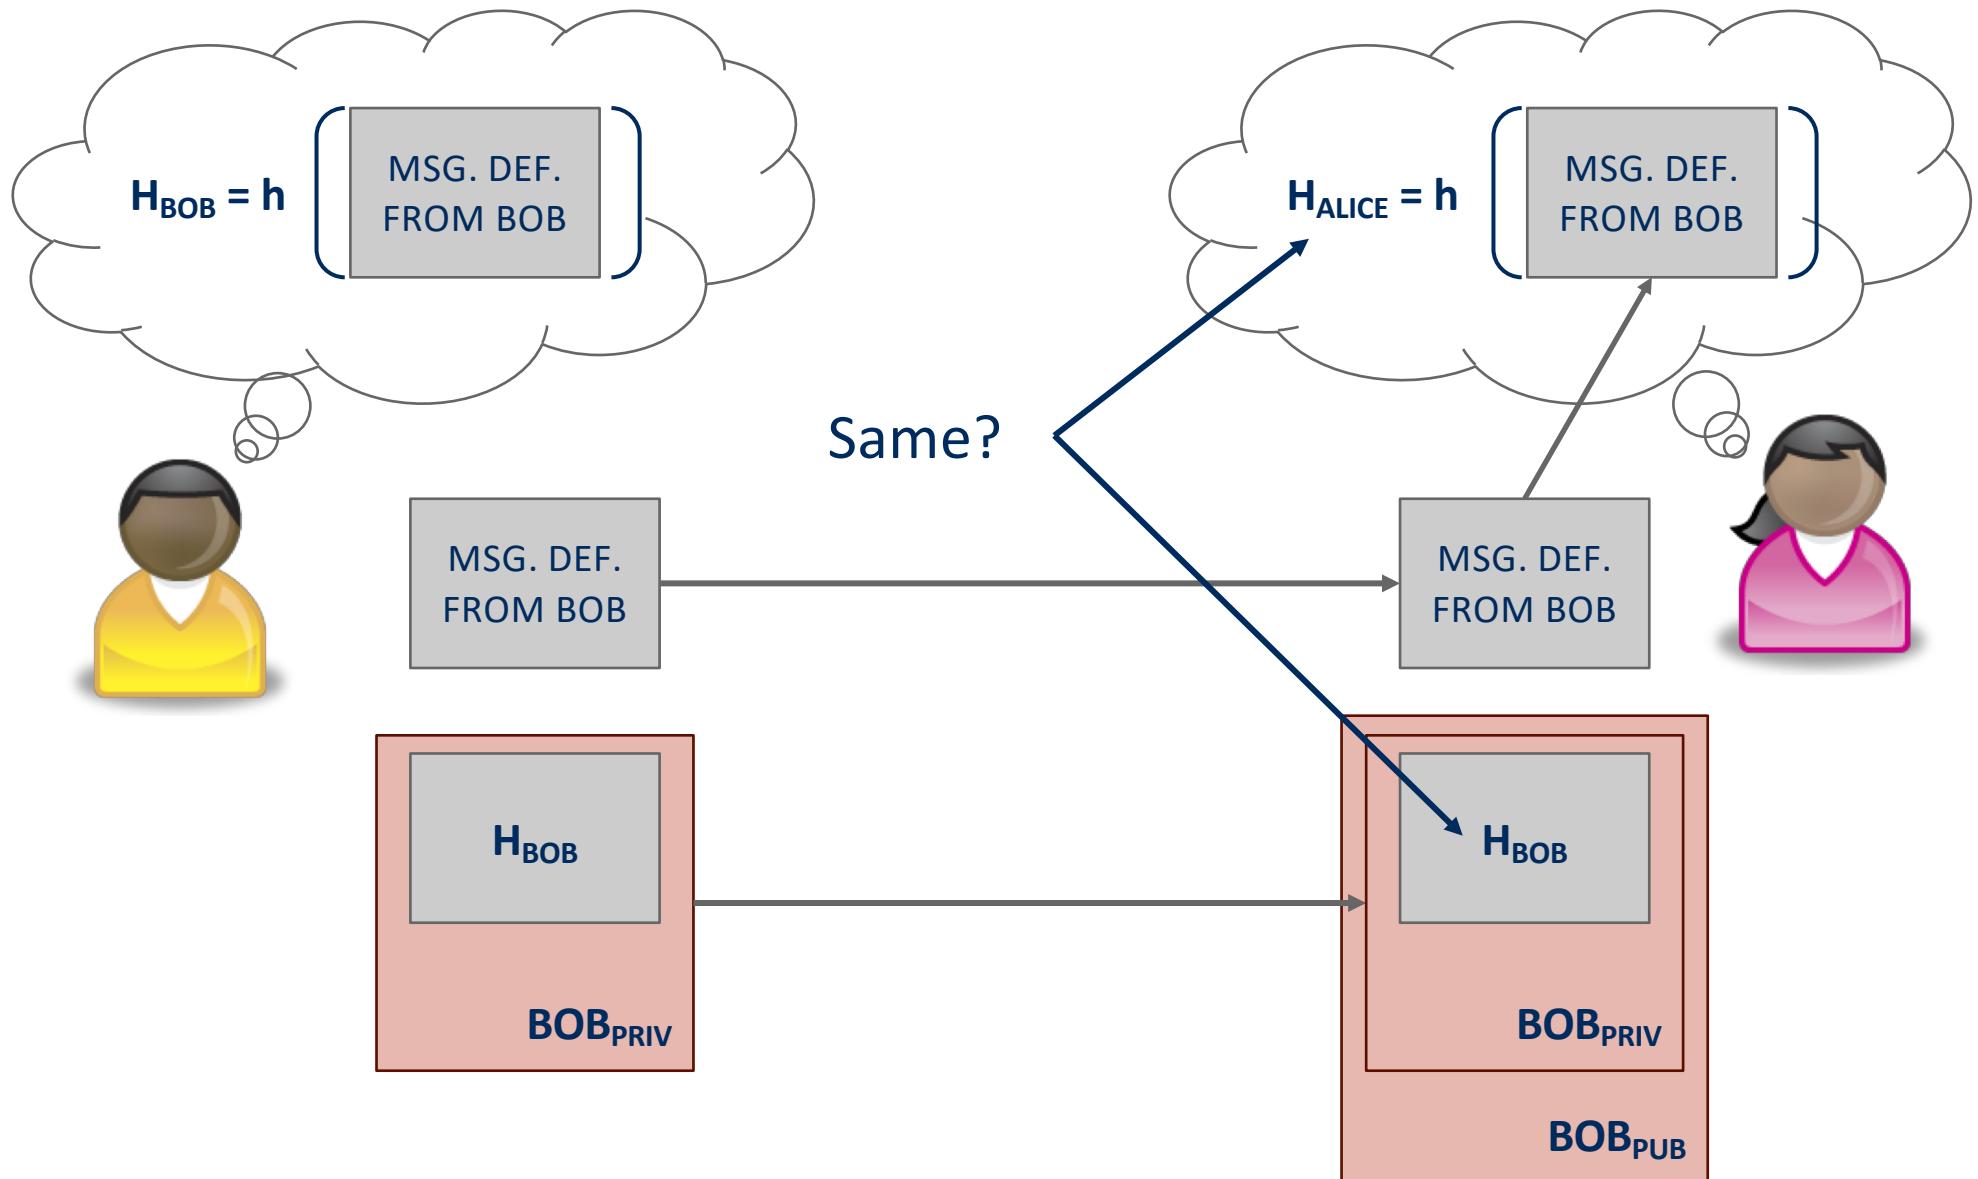
Do RSA signatures need to be slow?

- We encrypted symmetric crypto keys before to speed things up...
 - We'll need another crypto primitive to help out here
 - Cryptographically secure hash functions

Hashing for security (similarities)

- Cryptographically secure hash functions share properties with the hash functions we've already talked about:
 - Map from some input domain to a limited output range
 - Though output ranges are much larger here
 - For modern algorithms 224-512 bit output sizes
 - Time required to compute the hash is proportional to the size of the item being hashed
 - Though, practically, cryptographic hash functions are more expensive

Now just sign a hash of the message!



What about collisions?

- If Bob signs a hash of the message “I’ll see you at 7”
- It could appear that Bob signed any message whose hash collides with “I’ll see you at 7”...
- If $h(\text{"I'll see you at 7"}) == h(\text{"I'll see you after I rob the bank"})$, Bob could be in a lot of trouble
- An attack like this helped the Flame malware to spread
- This is also the reason Google is aiming to deprecate SHA-1

Hashing for security (differences)

- This is why cryptographically secure hash functions must support additional properties:
 - It should be infeasible to find two different messages with the same hash value
 - It should be infeasible to recover a message from its hash
 - Should require a brute force approach
 - Small changes to a message should change the hash value so extensively that the new hash value appears uncorrelated with the old hash value

Public key isn't perfect

What do you when a private key is compromised?

Final note about crypto

NEVER IMPLEMENT YOUR OWN CRYPTO

Use a trusted and tested library.